

Intelligent Systems in Medical Imaging

(NWI-IMC037)

Digital Image Processing

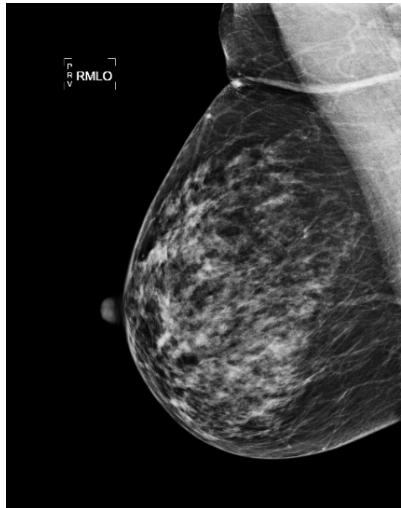
Francesco Ciompi

francesco.ciompi@radboudumc.nl

Radboud **umc**

Medical Imaging

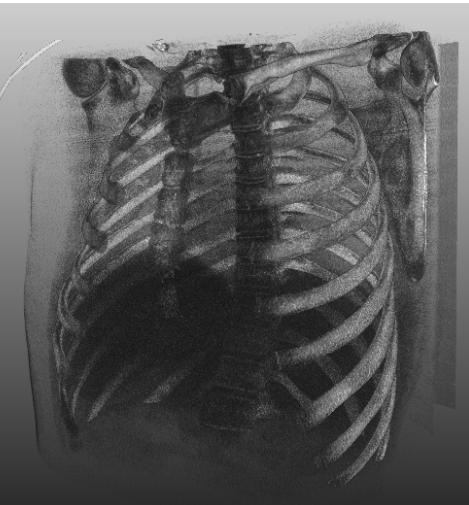
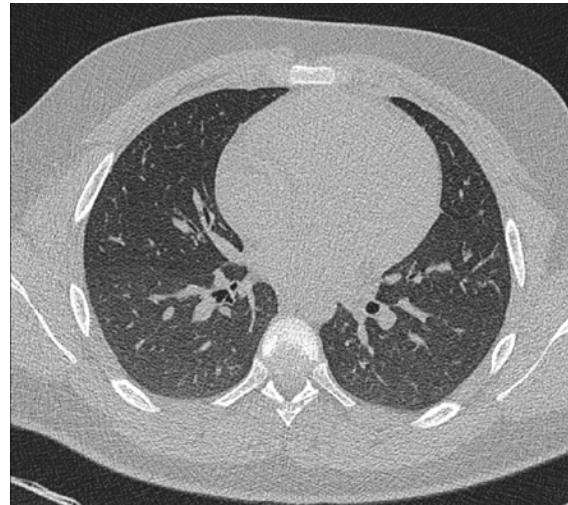
Mammography



2D, x-ray



Chest Computed Tomography (CT)

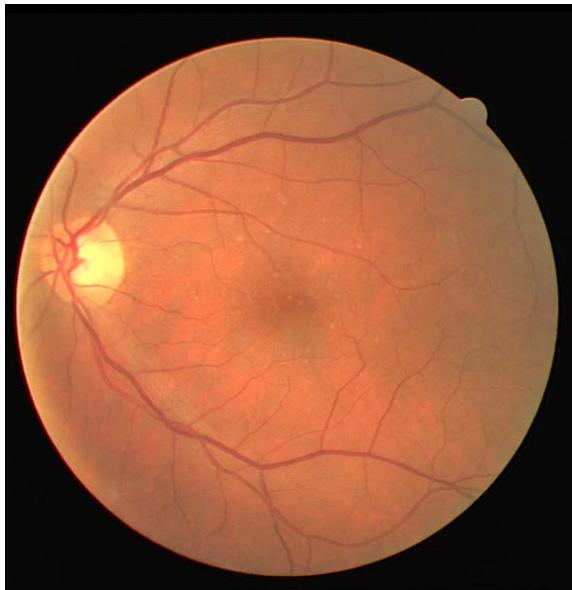


3D, x-ray



Medical Imaging

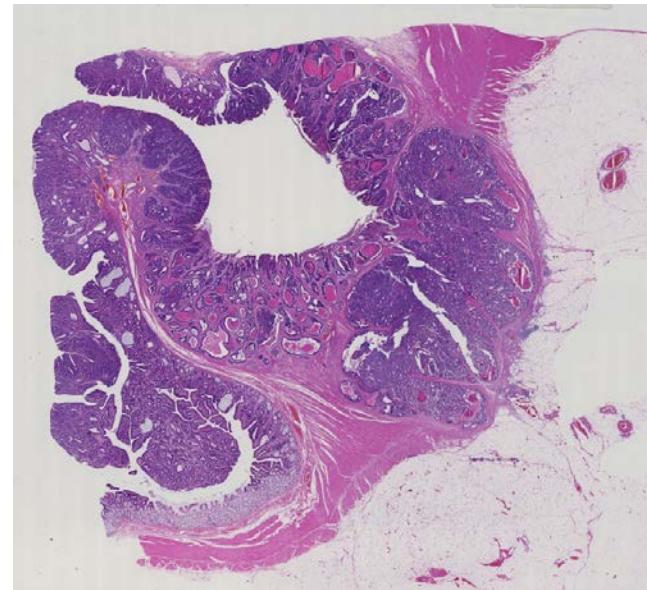
Retina fundus



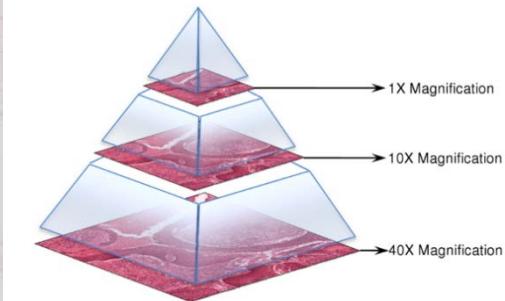
2D, RGB



Digital pathology

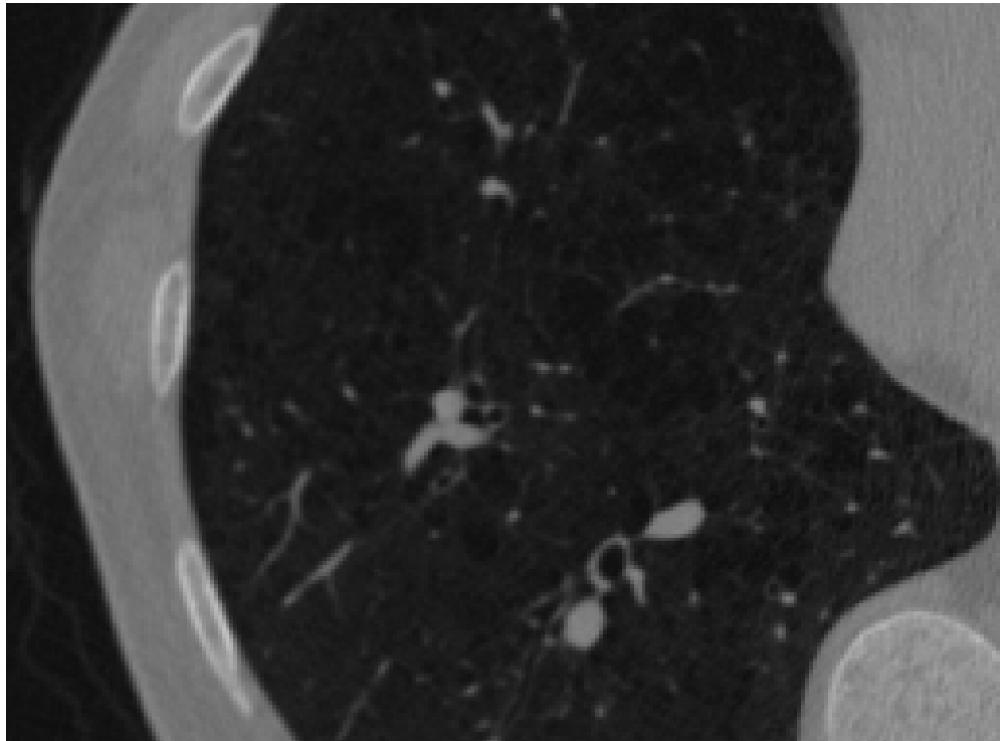


multi-resolution, RGB



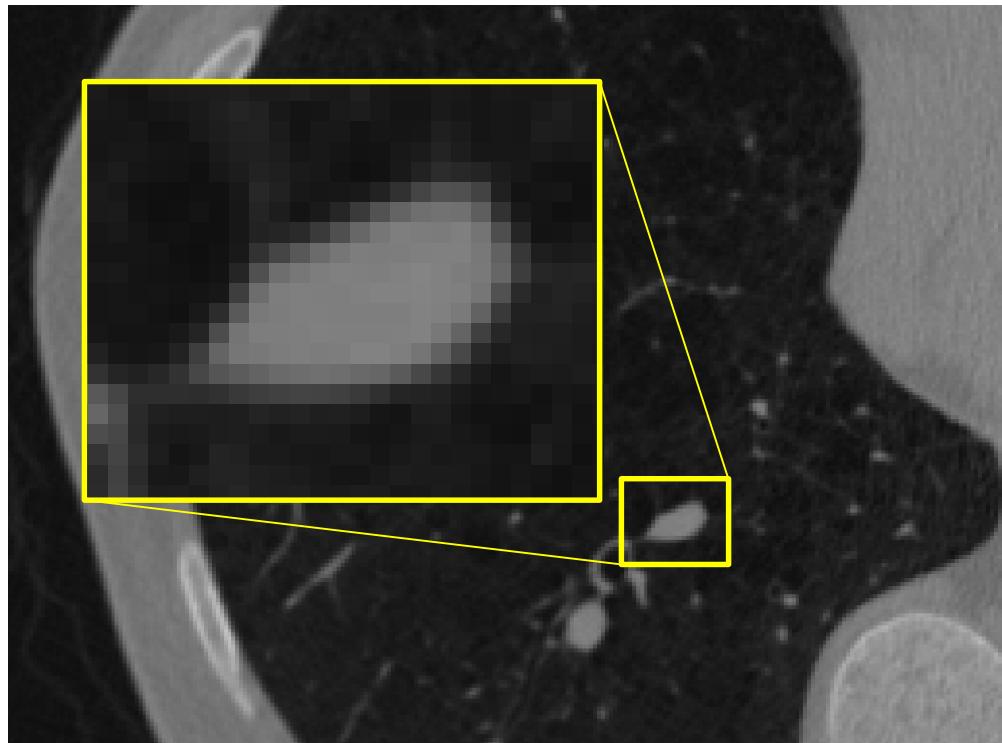
Digital image

- A digital image is a n-dimensional matrix of intensity values
- Each cell of the matrix is called pixel (or voxel)



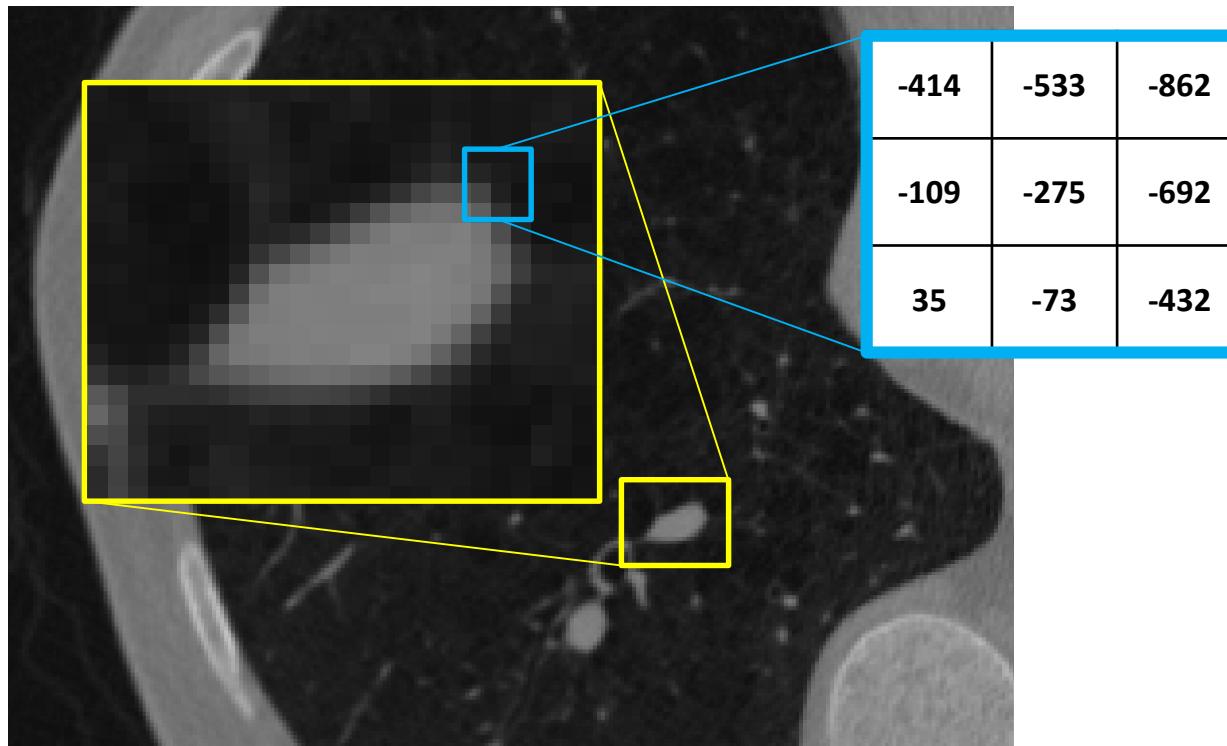
Digital image

- A digital image is a n-dimensional matrix of intensity values
- Each cell of the matrix is called pixel (or voxel)



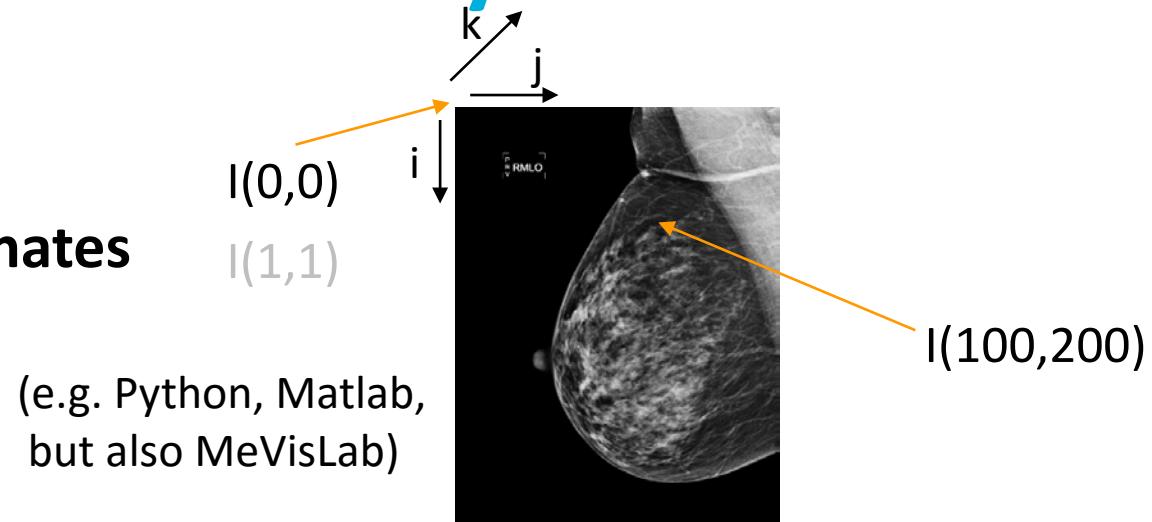
Digital image

- A digital image is a n-dimensional matrix of intensity values
- Each cell of the matrix is called pixel (or voxel)



Images and coordinate systems

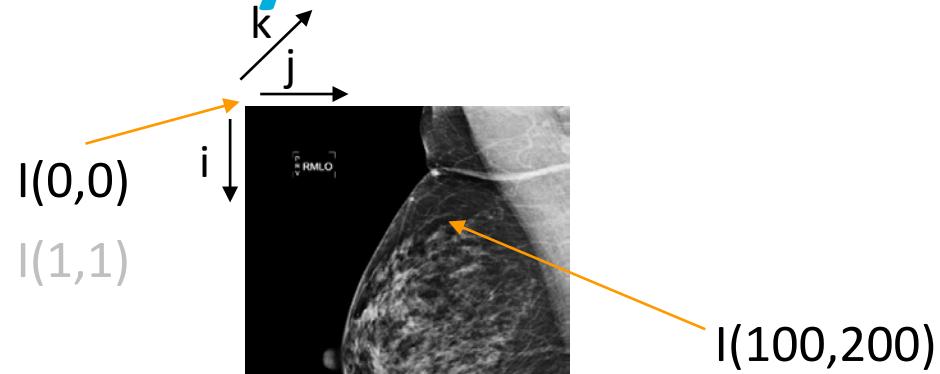
- Image indices (i, j, k)
 - **Pixel/voxel coordinates**



(e.g. Python, Matlab,
but also MeVisLab)

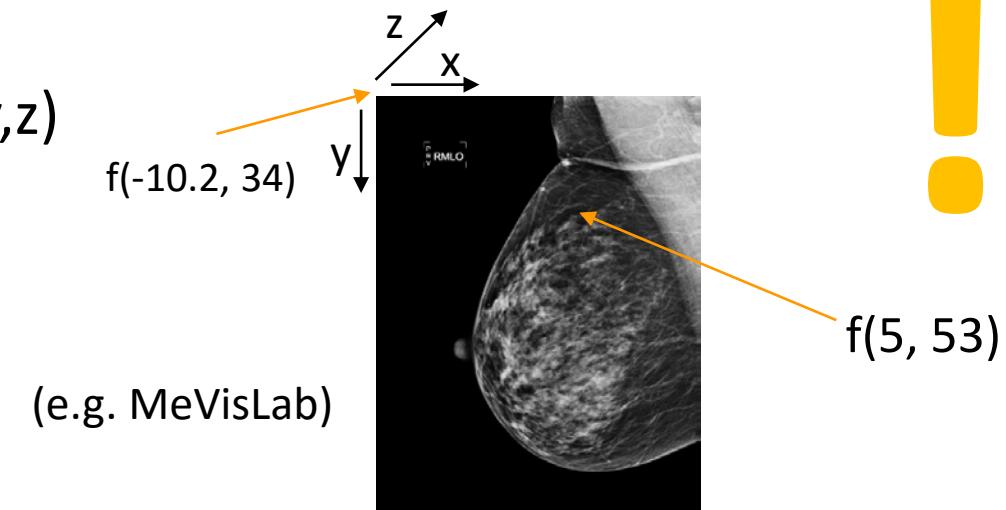
Images and coordinate systems

- Image indices (i, j, k)
 - **Pixel/voxel coordinates**

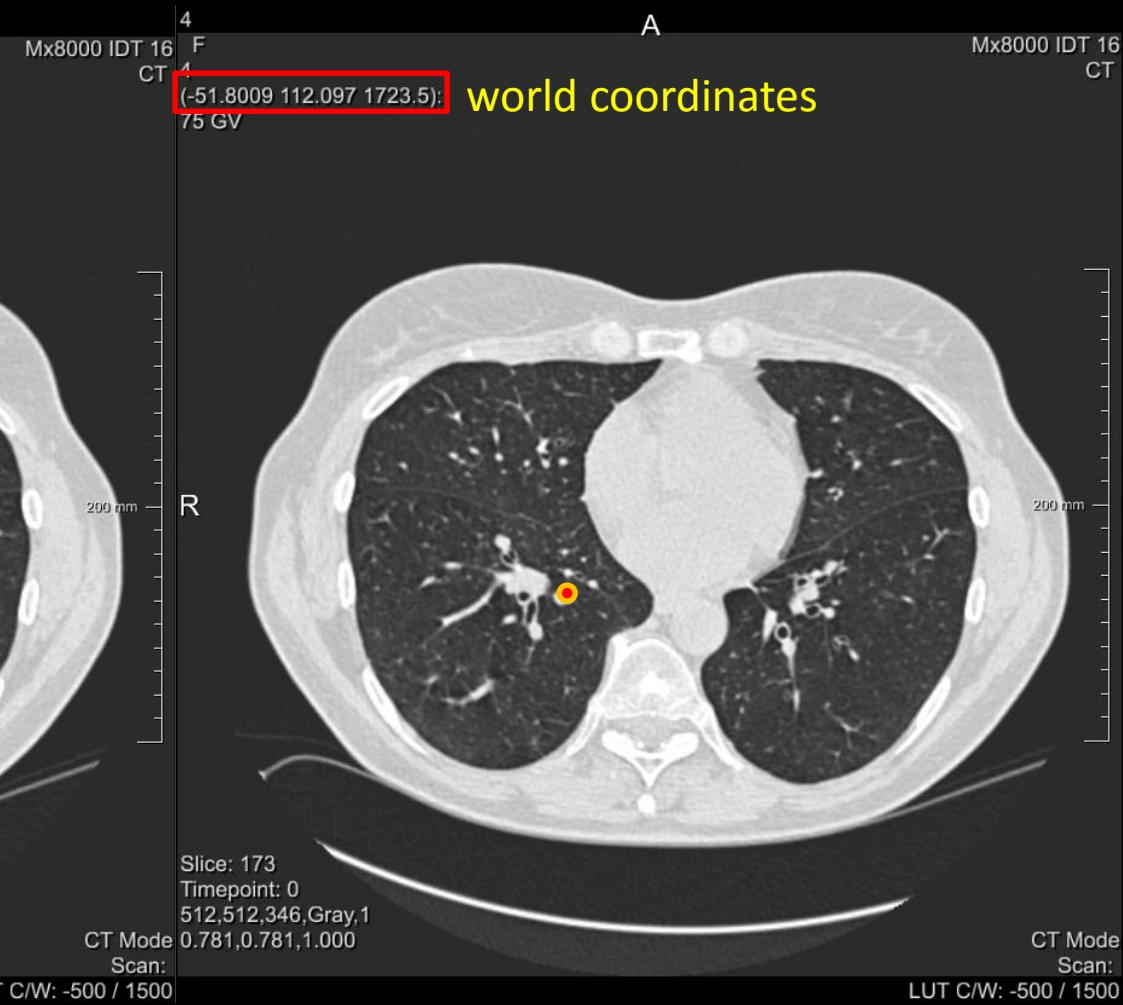
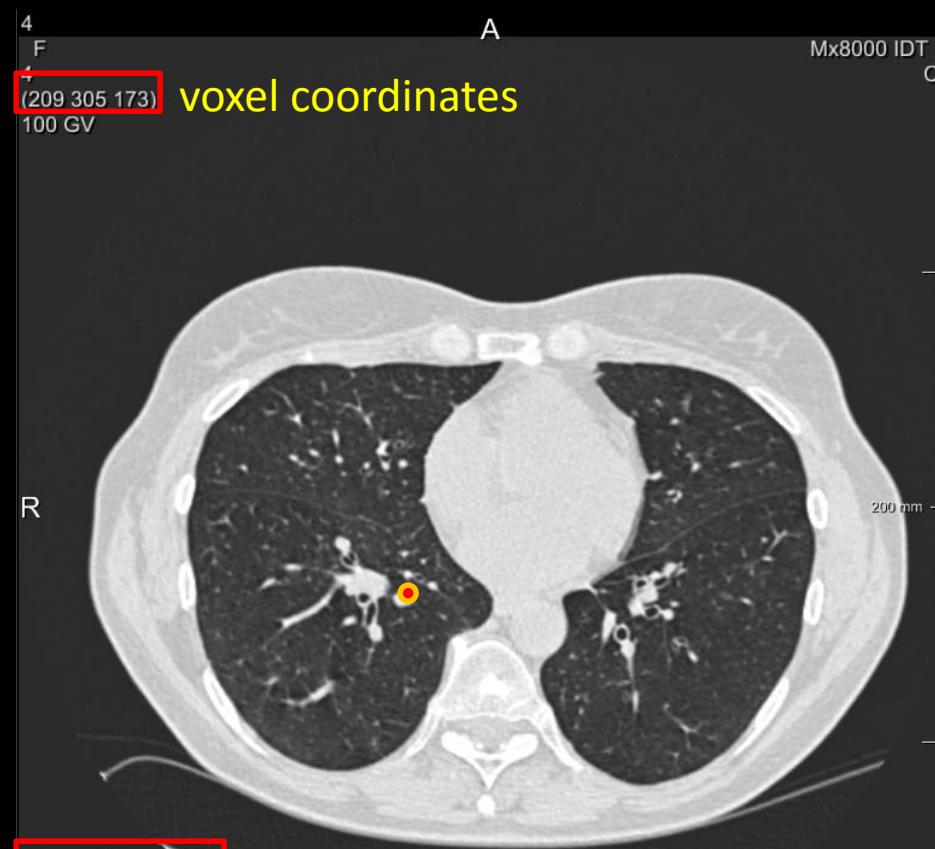


(e.g. Python, Matlab,
but also MeVisLab)

- Physical coordinates (x, y, z)
 - **World coordinates**
 - Origin
 - Orientation
 - Spacing

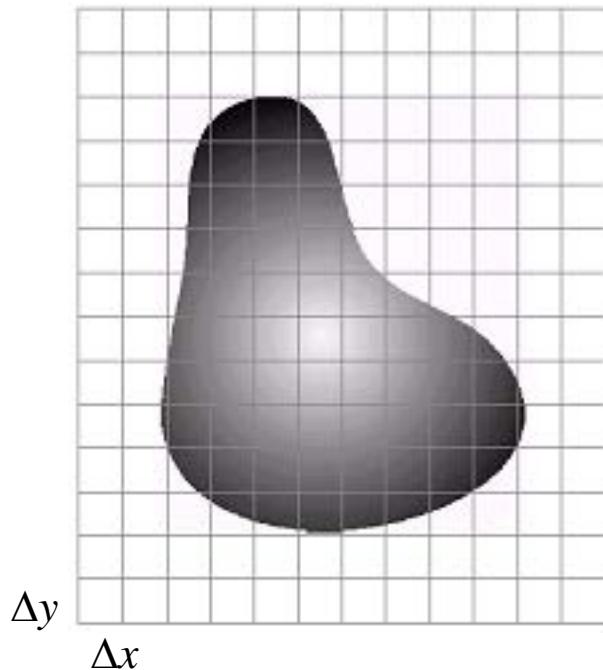


(e.g. MeVisLab)

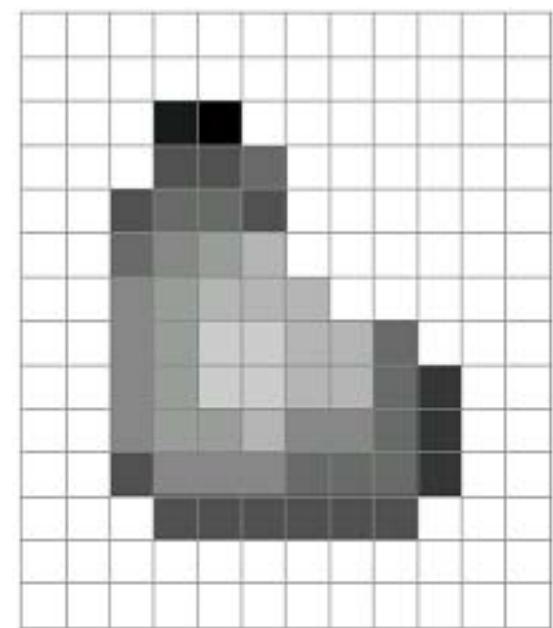


Sampling and quantization

We have to convert data from the real world into digital images



SAMPLING
+
QUANTIZATION

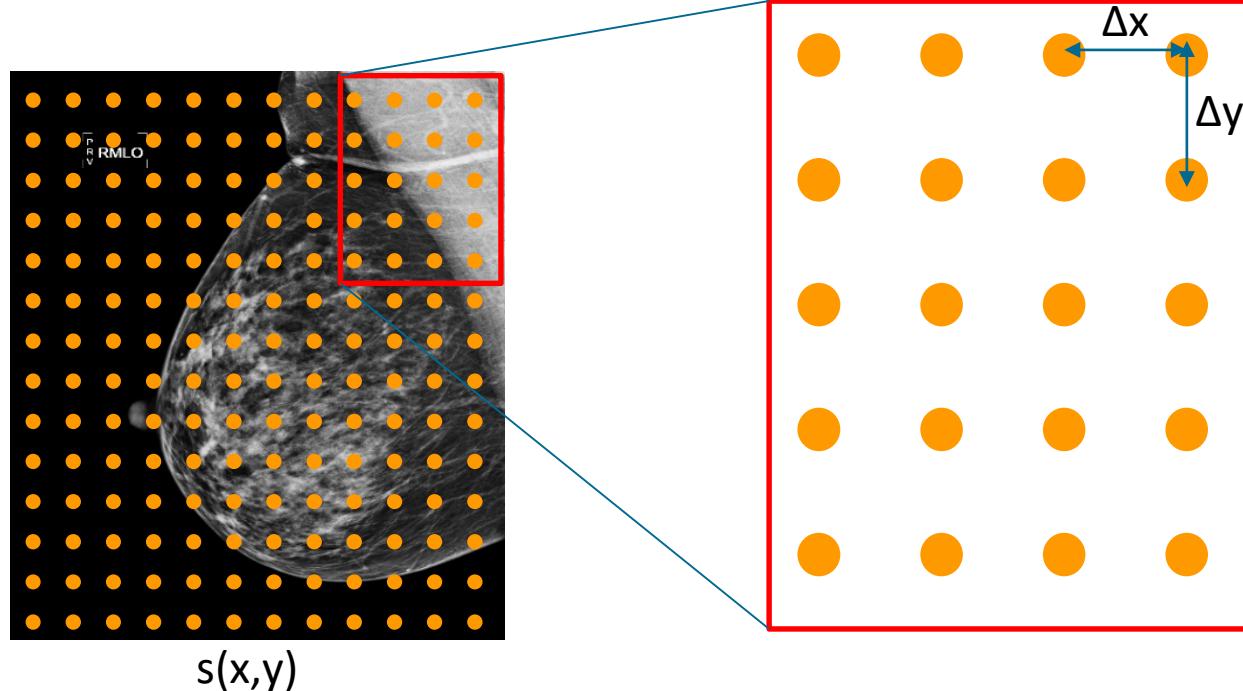


Sampling

The ideal sampling $s(x,y)$ in a regular grid can be represented using a collection of Dirac functions

$$s(x, y) = \sum_{j=1}^M \sum_{k=1}^N \delta(x - j\Delta x, y - k\Delta y)$$

SPACING



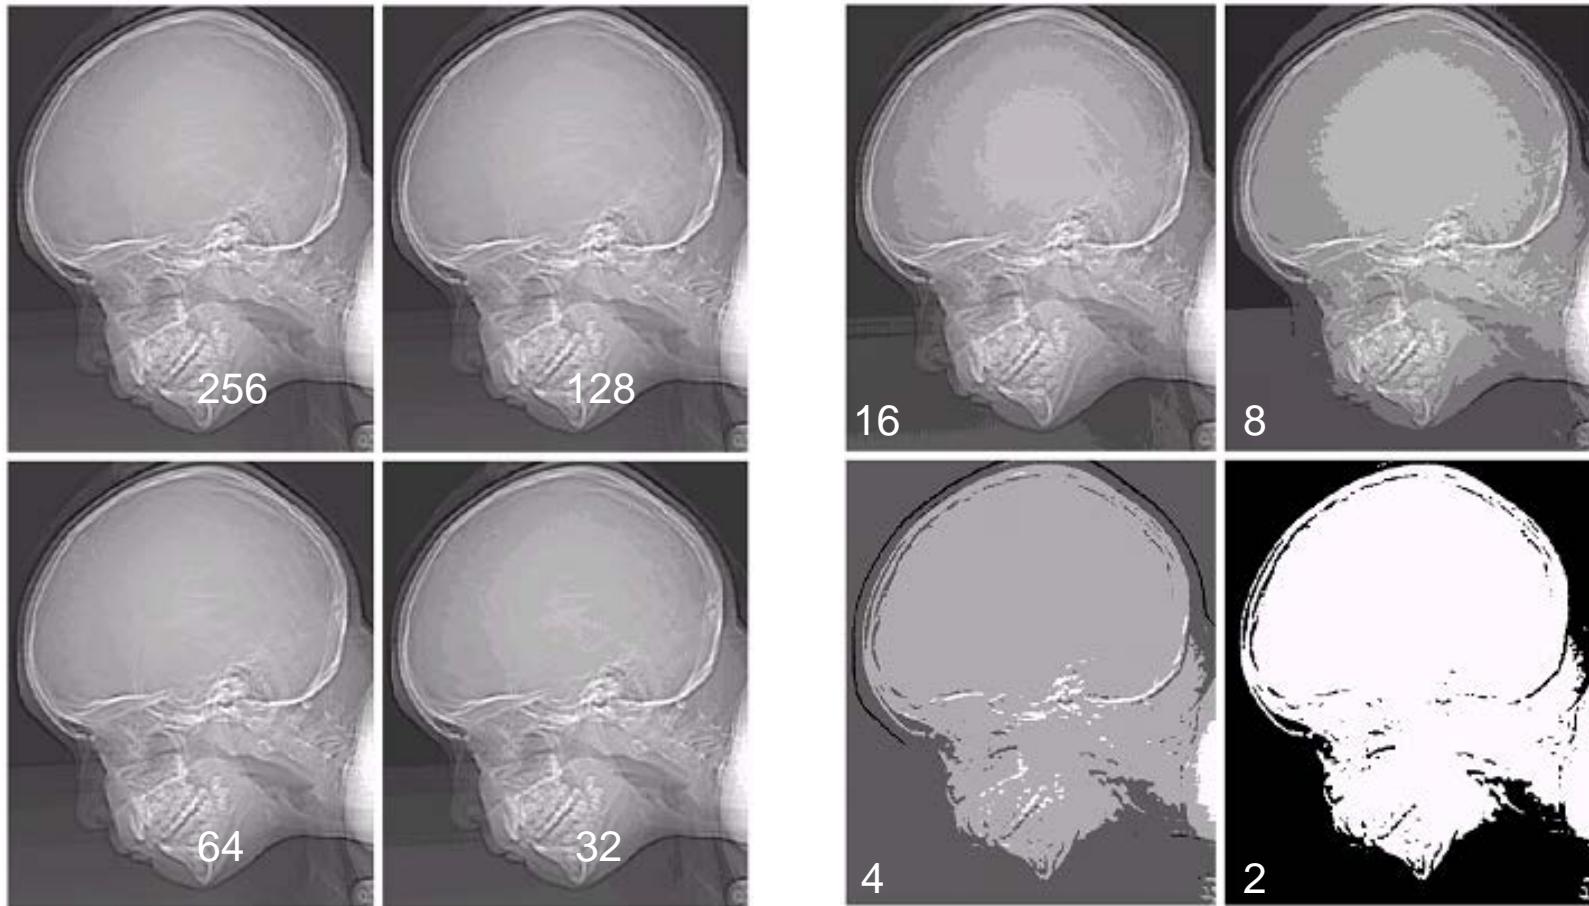
DICOM header

(0008, 0032) Acquisition Time
(0008, 0033) Content Time
(0008, 0050) Accession Number
(0008, 0060) Modality
(0008, 0070) Manufacturer
(0008, 0090) Referring Physician's Name
(0008, 1090) Manufacturer's Model Name
(0008, 1155) Referenced SOP Instance UID
(0010, 0010) Patient's Name
(0010, 0020) Patient ID
(0010, 0030) Patient's Birth Date
(0010, 0040) Patient's Sex
(0010, 1010) Patient's Age
(0010, 2100) Last Menstrual Date
(0012, 0062) Patient Identity Removed
(0012, 0063) De-Identification Method
(0013, 0010) Private Creator
(0013, 1010) Private tag data
(0013, 1013) Private tag data
(0018, 0010) Contrast/Bolus Agent
(0018, 0015) Body Part Examined
(0018, 0022) Scan Options
(0018, 0050) Slice Thickness
(0018, 0060) KVP
(0018, 0090) Data Collection Diameter
(0018, 1020) Software Version(s)
(0018, 1100) Reconstruction Diameter
(0018, 1110) Distance Source to Detector
(0018, 1111) Distance Source to Patient
(0018, 1120) Gantry/Detector Tilt
(0018, 1130) Table Height
(0018, 1140) Rotation Direction
(0018, 1150) Exposure Time
(0018, 1151) X-Ray Tube Current
(0018, 1152) Exposure
(0018, 1160) Filter Type
(0018, 1170) Generator Power
(0018, 1190) Focal Spot(s)
(0018, 1210) Convolution Kernel
(0018, 5100) Patient Position
(0020, 000d) Study Instance UID
(0020, 000e) Series Instance UID
(0020, 0010) Study ID
(0020, 0011) Series Number
(0020, 0013) Instance Number
(0020, 0032) Image Position (Patient)
(0020, 0037) Image Orientation (Patient)
(0020, 0052) Frame of Reference UID
(0020, 1040) Position Reference Indicator
(0020, 1041) Slice Location
(0028, 0002) Samples per Pixel
(0028, 0004) Photometric Interpretation
(0028, 0010) Rows
(0028, 0011) Columns
(0028, 0030) Pixel Spacing
(0028, 0100) Bits Allocated
(0028, 0101) Bits Stored
(0028, 0102) High Bit
(0028, 0103) Pixel Representation
(0028, 0120) Pixel Padding Value
(0028, 0303) Longitudinal Temporal Information M CS: 'MODIFIED'
(0028, 1050) Window Center
 DS: '-600'
 TM: ''
 TM: ''
 SH: '2819497684894126'
 CS: 'CT'
 LO: 'GE MEDICAL SYSTEMS'
 PN: ''
 LO: 'LightSpeed Plus'
 UI: 1.3.6.1.4.1.14519.5.2.1.6279.6001.675906998158803995297223798692
 PN: ''
 LO: 'LIDC-IDRI-0001'
 DA: ''
 CS: ''
 AS: ''
 DA: '20000101'
 CS: 'YES'
 LO: 'DCM:113100/113105/113107/113108/113109/113111'
 LO: 'CTP'
 LO: 'LIDC-IDRI'
 LO: '62796081'
 LO: 'IV'
 CS: 'CHEST'
 CS: 'HELICAL MODE'
 DS: '2.500000'
 DS: '120'
 DS: '500.00000'
 LO: 'LightSpeed4Apps2.4.2_H2.4M5'
 DS: '360.00000'
 DS: '949.075012'
 DS: '541.00000'
 DS: '0.000000'
 DS: '144.399994'
 CS: 'CW'
 IS: '570'
 IS: '400'
 IS: '4684'
 SH: 'BODY FILTER'
 IS: '48000'
 DS: '1.200000'
 SH: 'STANDARD'
 CS: 'FFS'
 UI: 1.3.6.1.4.1.14519.5.2.1.6279.6001.298806137288633453246975630178
 UI: 1.3.6.1.4.1.14519.5.2.1.6279.6001.179049373636438705059720603192
 SH: ''
 IS: '3000566'
 IS: '80'
 DS: '[-166.000000, -171.699997, -207.500000]'
 DS: '[1.000000, 0.000000, 0.000000, 1.000000, 0.000000]'
 UI: 1.3.6.1.4.1.14519.5.2.1.6279.6001.229925374658226729607867499499
 LO: 'SN'
 DS: '-207.500000'
 US: 1
 CS: 'MONOCHROME2'
 US: 512
 US: 512
 DS: '1.0.703125' '0.703125'
 US: 16
 US: 16
 US: 15
 US: 1
 US: 63536
 CS: 'MODIFIED'
 DS: '-600'

Pixel spacing

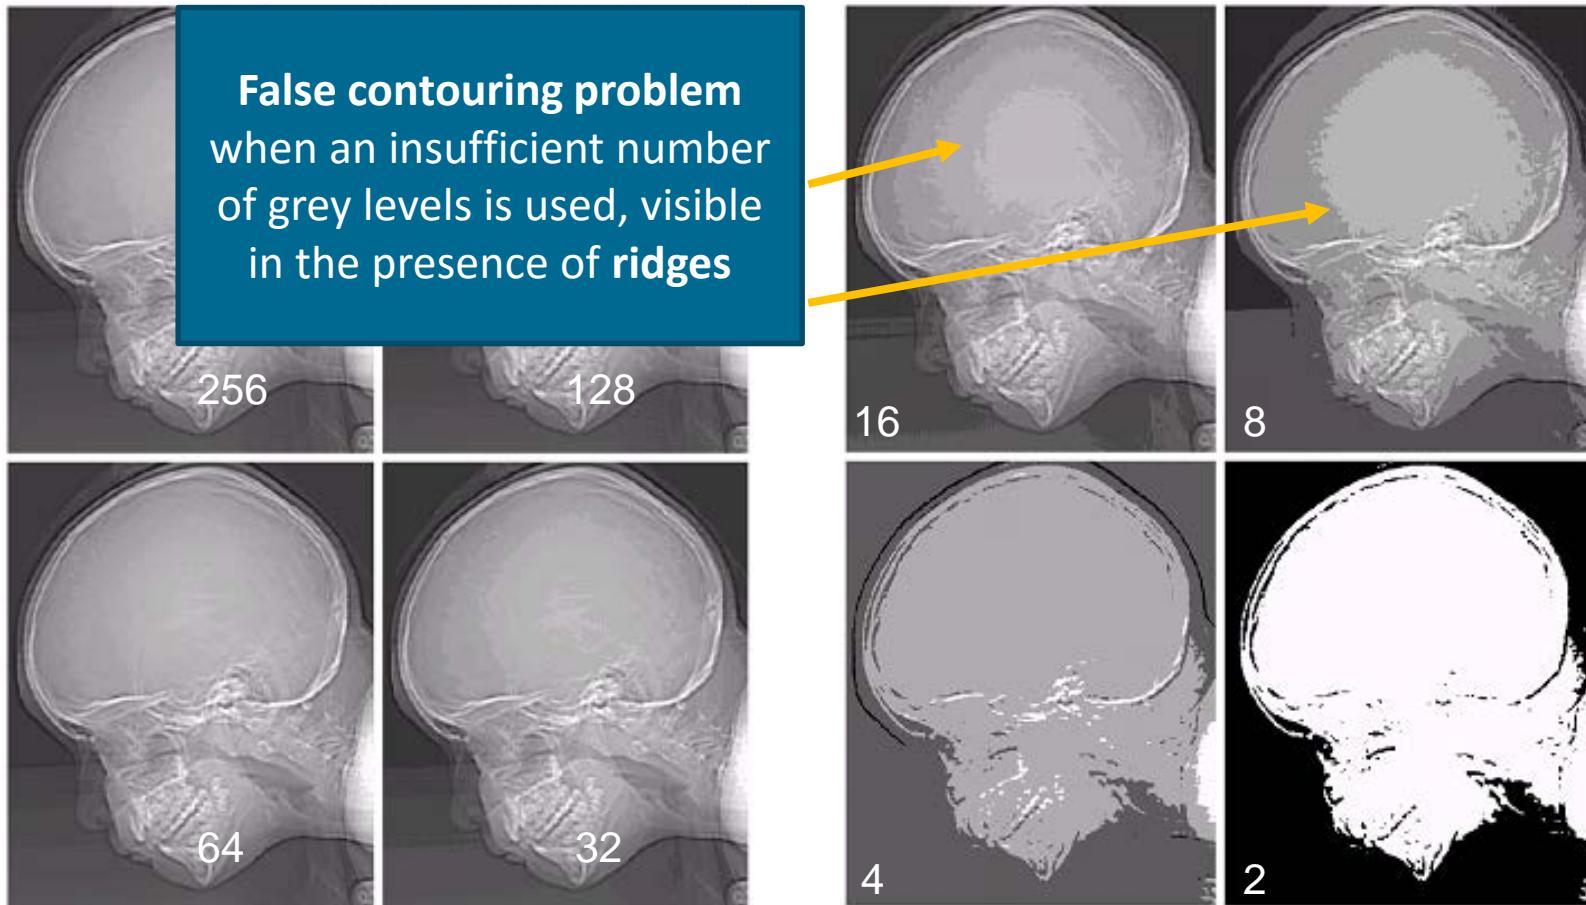
Quantization

X-ray image



Quantization

X-ray image



Quantization

In **greyscale** digital images, the intensity is quantized using 256 levels: [0-255]

In **color** digital images, three channels (R,G,B) with intensity quantized with 256 levels are used: [[0-255], [0-255], [0-255]]

In **X-Ray Computed Tomography**,
the intensity is measured in
Hounsfield Units (HU), in a range
[-1000,1000] (3000 for dense bone)

Tissue	CT Number (HU)
Bone	+1000
Liver	40-60
White mater	-20 to -30
Grey mater	-37 to -45
Blood	40
Muscle	10-40
Kidney	30
CSF	15
Water	0
Fat	-50 to -100
Air	-1000

Intensity transformation

Intensity transformation

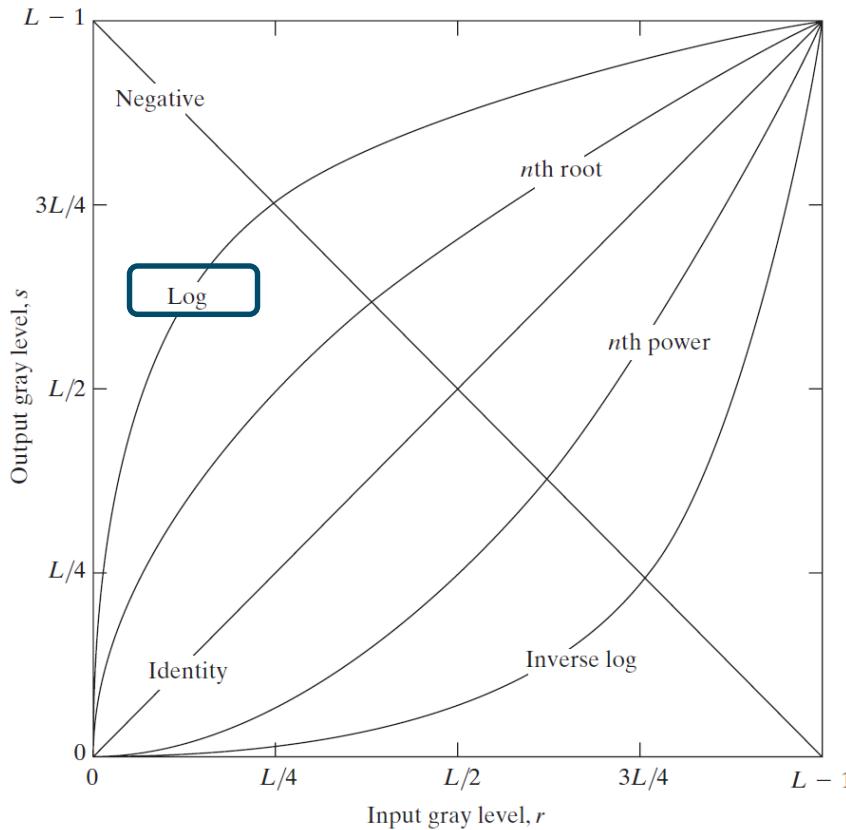
A transformation T of the original pixel value p from scale $[p_0, p_k]$ into brightness from a new scale $[q_0, q_k]$ is given by:

$$q = T(p)$$

The correction depends on the pixel **value**,
not on its position (neighborhood)

Intensity transformation functions

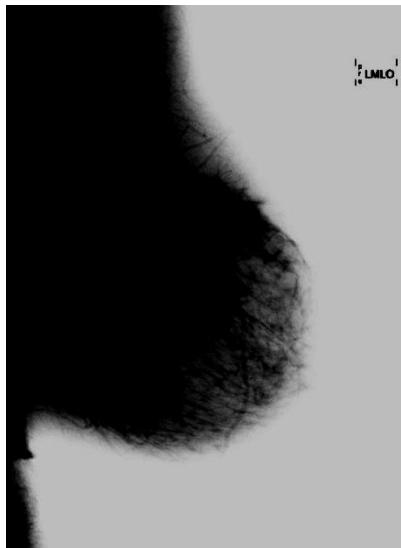
Some common intensity transformation functions:



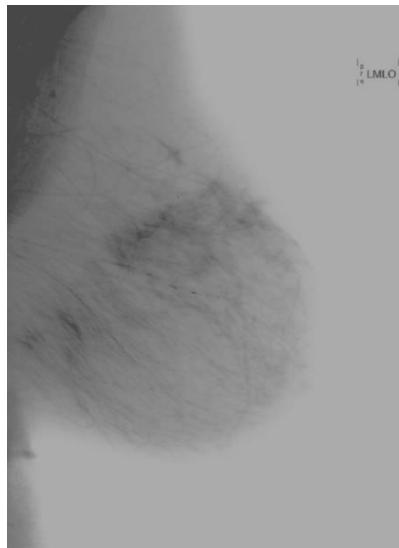
- In natural images, $L = 256$
- In medical images, it depends on the image modality
- The **Log** transformation **boosts** low grey-level values
- The “negative” function is also called “invert” function

Example: Mammogram processing

RAW



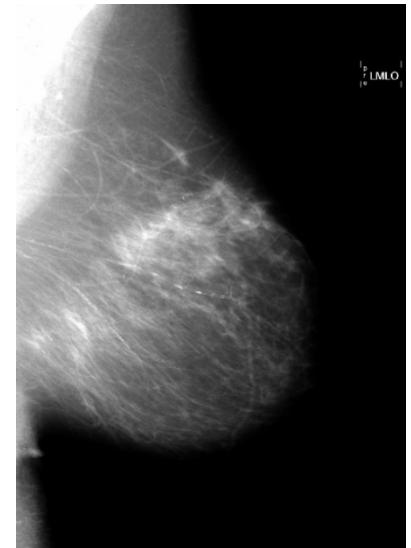
LOG



INVERT



CONTRAST STRETCH

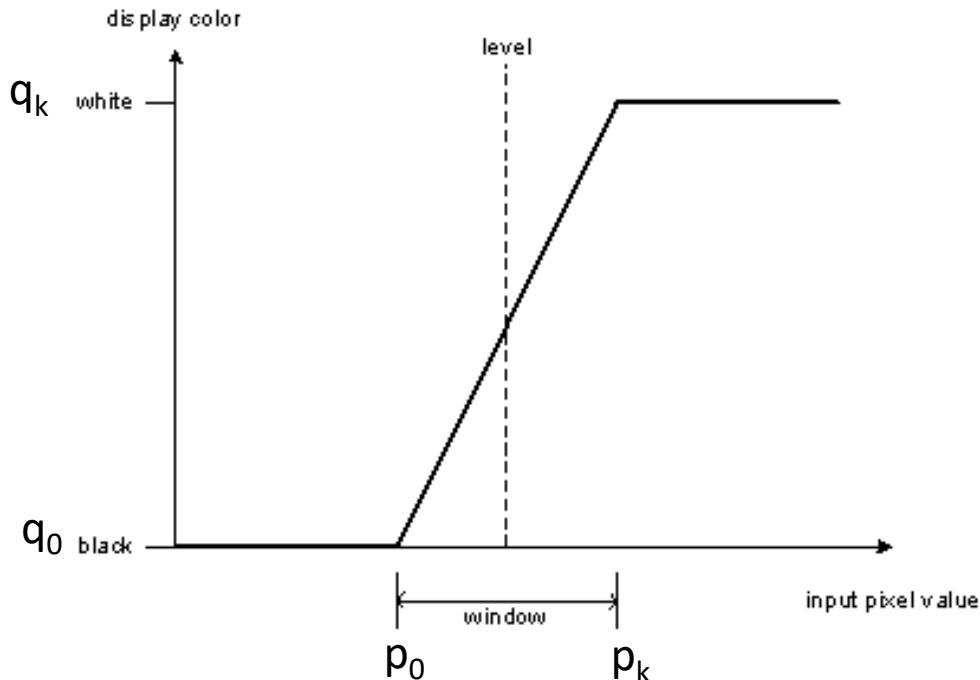


$$I(x, y) = I_0 e^{-\mu(x, y)d(x, y)}$$

$$\log(I) = \log(I_0) - \mu(x, y)d(x, y)$$



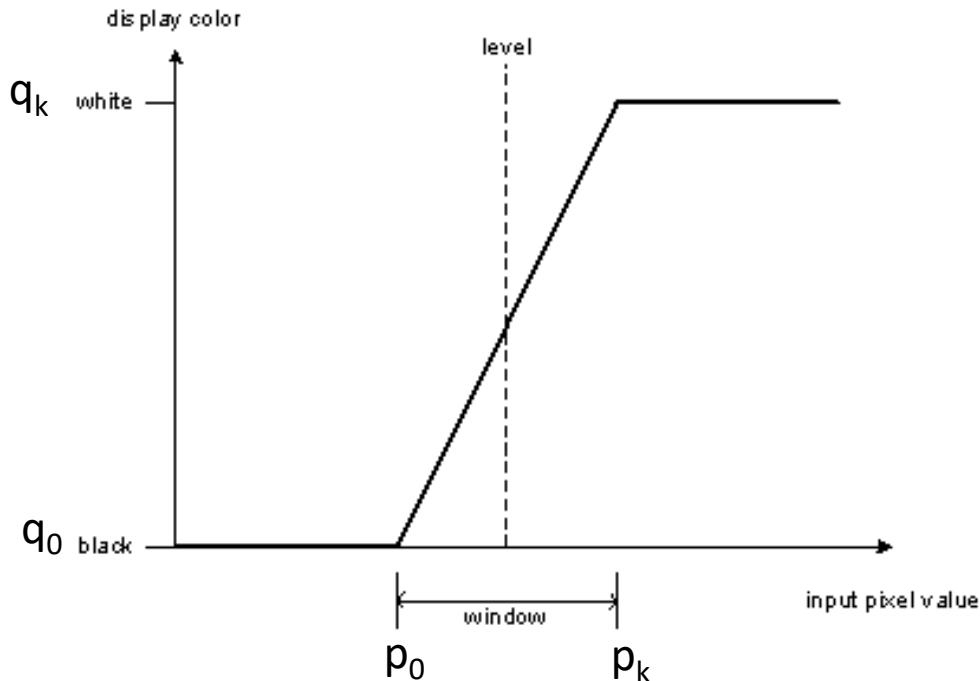
Linear contrast stretching



- map the range $[p_0, p_k]$ to the range $[q_0, q_k]$
- first rescale the range $[p_0, p_k]$ to $[0, 1]$
- then rescale the result to the desired range $[q_0, q_k]$

$$q = T(p) = q_0 + (q_k - q_0) \left[\frac{p - p_0}{p_k - p_0} \right]$$

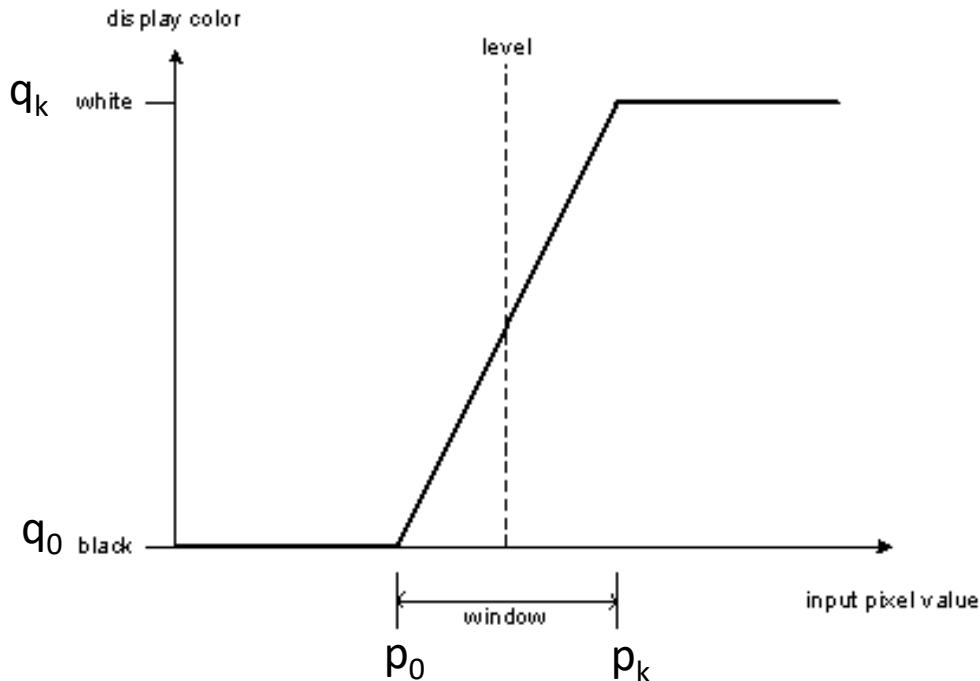
Linear contrast stretching



- map the range $[p_0, p_k]$ to the range $[q_0, q_k]$
- first rescale the range $[p_0, p_k]$ to $[0, 1]$
- then rescale the result to the desired range $[q_0, q_k]$

$$q = T(p) = q_0 + (q_k - q_0) \left[\frac{p - p_0}{p_k - p_0} \right]$$

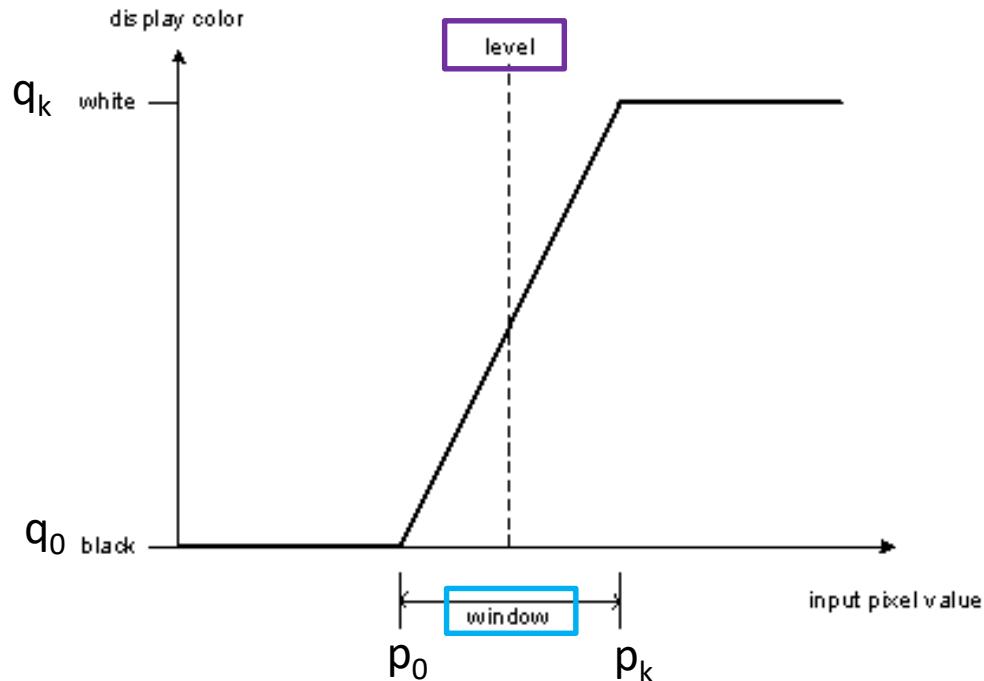
Linear contrast stretching



- map the range $[p_0, p_k]$ to the range $[q_0, q_k]$
- first rescale the range $[p_0, p_k]$ to $[0, 1]$
- then rescale the result to the desired range $[q_0, q_k]$

$$q = T(p) = \boxed{q_0 + (q_k - q_0) \left[\frac{p - p_0}{p_k - p_0} \right]}$$

Linear contrast stretching



- map the range $[p_0, p_k]$ to the range $[q_0, q_k]$
- first rescale the range $[p_0, p_k]$ to $[0, 1]$
- then rescale the result to the desired range $[q_0, q_k]$

$$q = T(p) = q_0 + (q_k - q_0) \left[\frac{p - p_0}{p_k - p_0} \right]$$

Linear contrast stretching

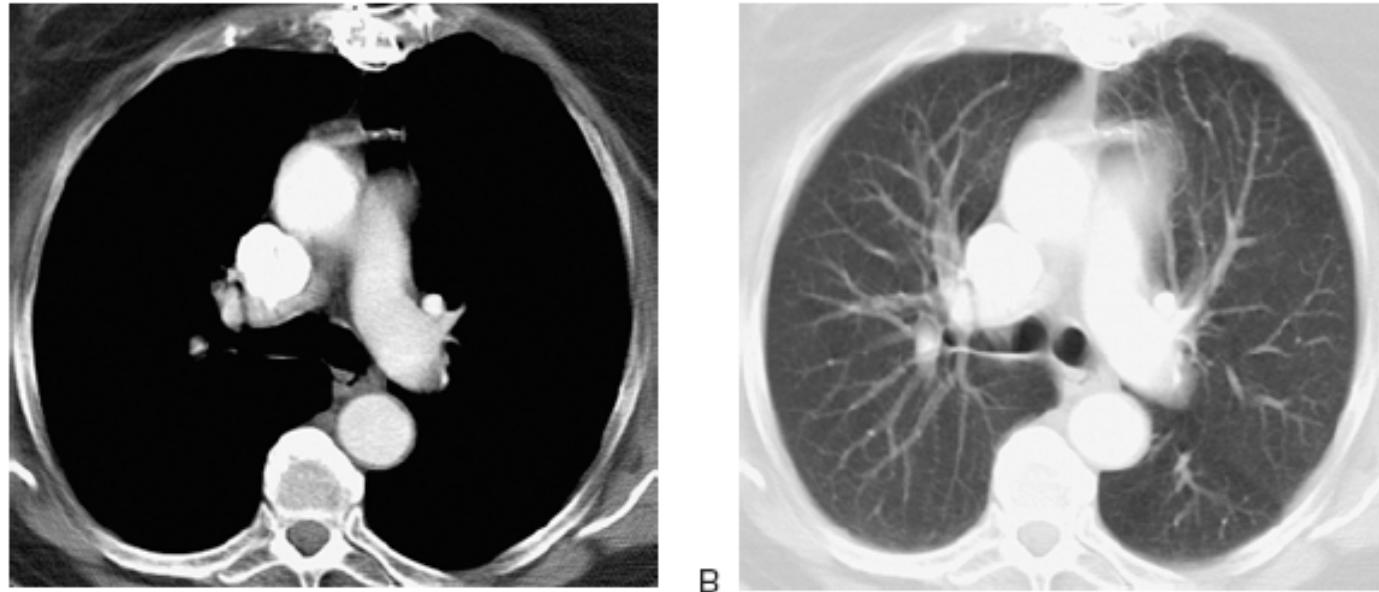
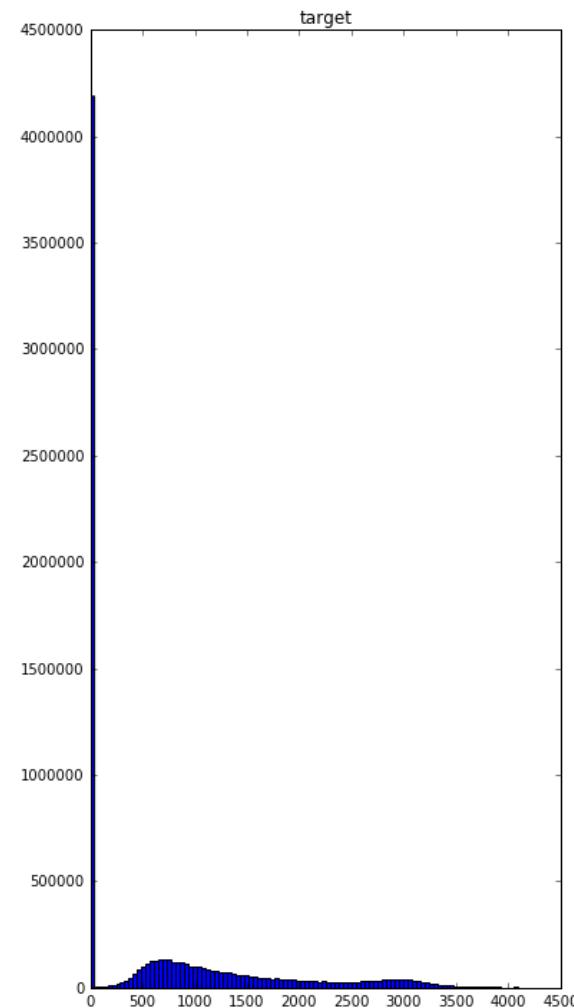
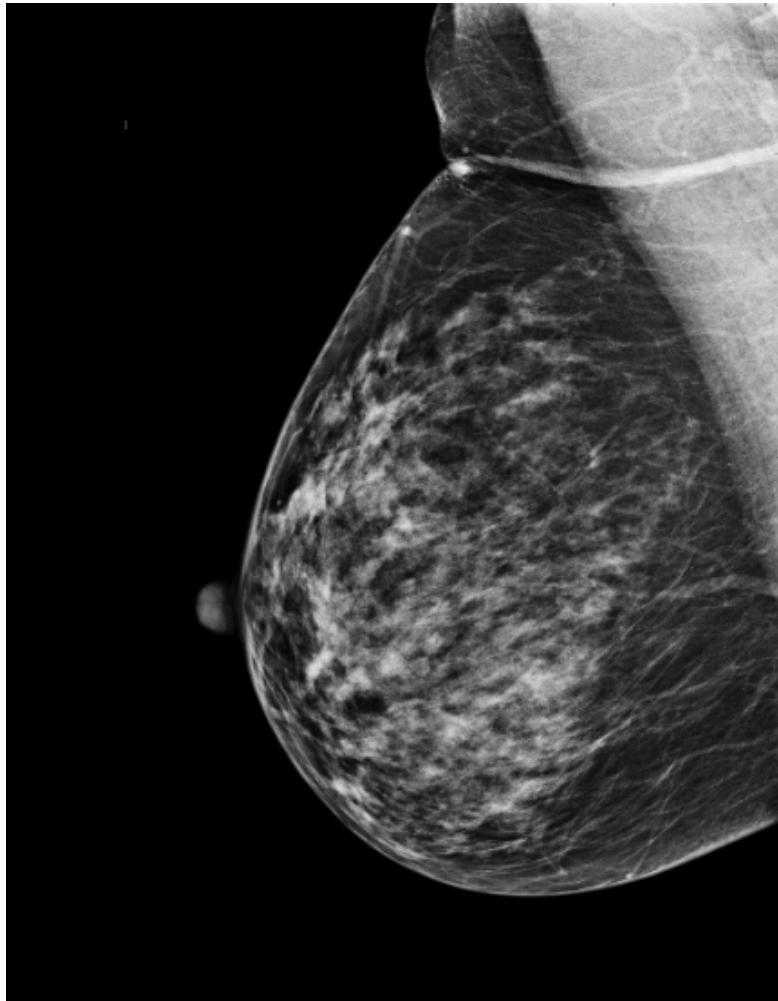
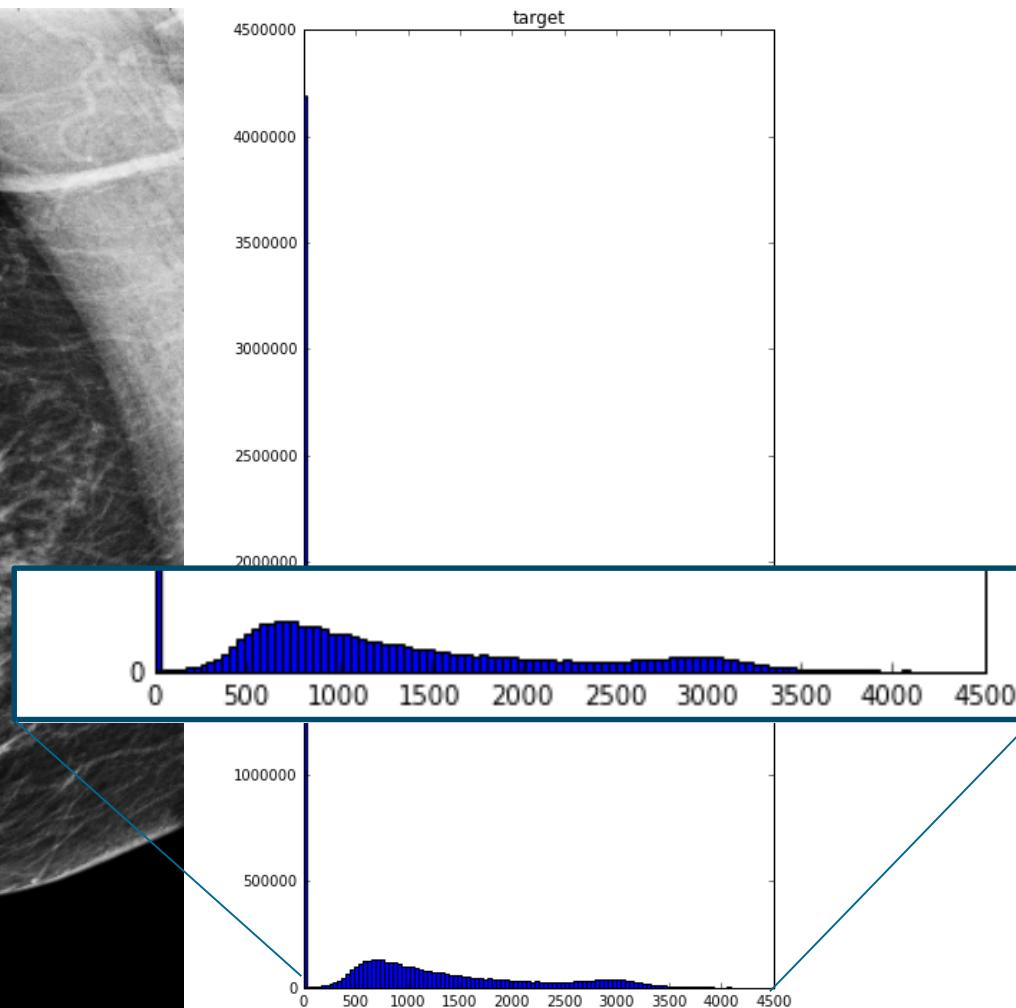
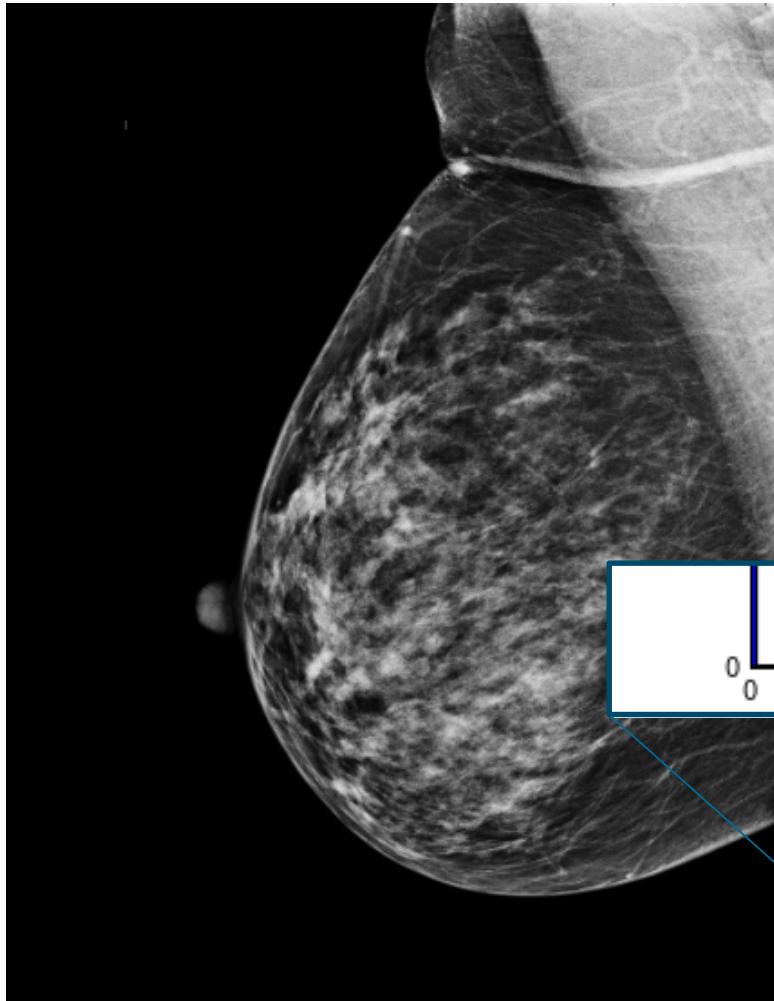


Fig. 1.5 These two images are of the same section, viewed at different window settings. **(A)** A window level of +40 with a window width of 350 reveals structures within the mediastinum but no lung parenchyma can be seen. **(B)** The window level is -600 with a window width of 1500 Hounsfield units. This enables details of the lung parenchyma to be seen, at the expense of the mediastinum.

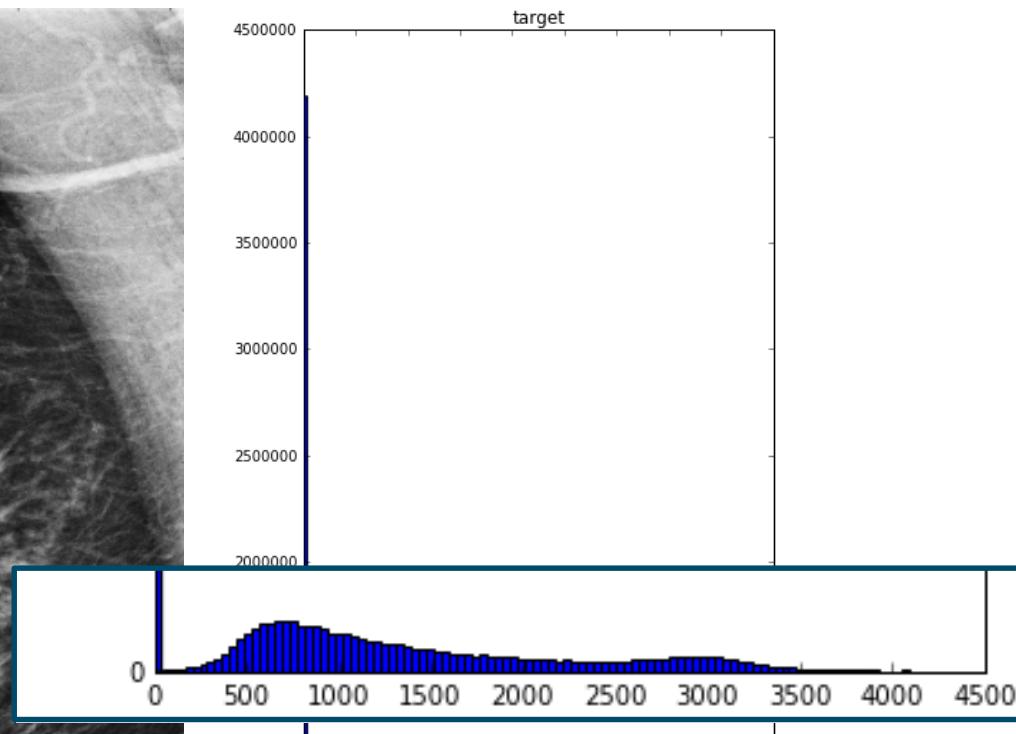
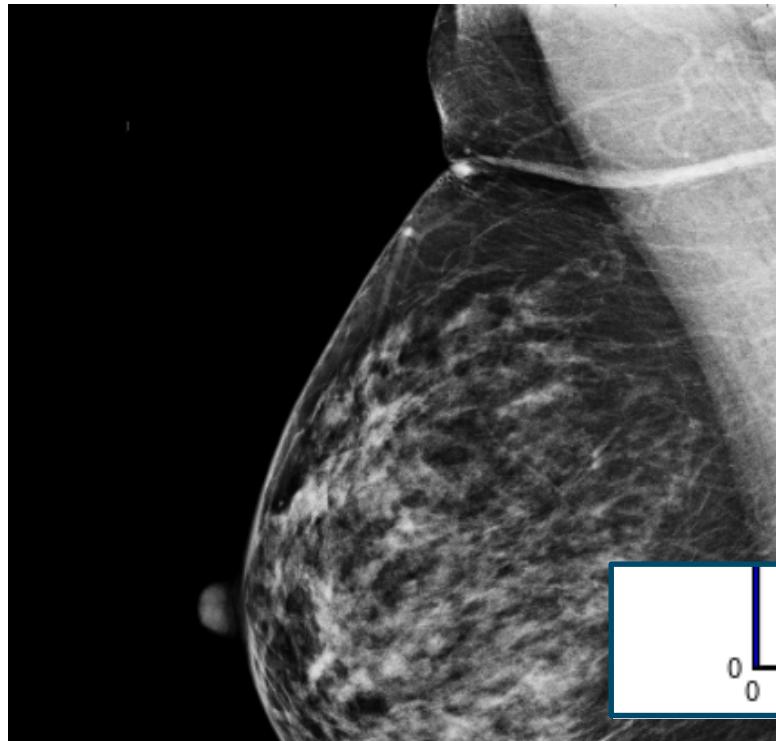
Histograms



Histograms



Histograms



The **histogram** $h_f(p)$ of an image provides the frequency of the pixel value p in the image. The **normalized histogram** is obtained by dividing by the number of pixels (voxels).

In Python



[Scipy.org](#) [Docs](#) [NumPy v1.12 Manual](#) [NumPy Reference](#) [Routines](#) [Statistics](#)

numpy.histogram

`numpy.histogram(a, bins=10, range=None, normed=False, weights=None, density=None)`

[\[source\]](#)

Compute the histogram of a set of data.

Parameters: `a : array_like`

Input data. The histogram is computed over the flattened array.

`bins : int or sequence of scalars or str, optional`

If `bins` is an int, it defines the number of equal-width bins in the given range (10, by default). If `bins` is a sequence, it defines the bin edges, including the rightmost edge, allowing for non-uniform bin widths.

New in version 1.11.0.

If `bins` is a string from the list below, `histogram` will use the method chosen to calculate the optimal bin width and consequently the number of bins (see *Notes* for more detail on the estimators) from the data that falls within the requested range. While the bin width will be optimal for the actual data in the range, the number of bins will be computed to fill the entire range, including the empty portions. For visualisation, using the 'auto' option is suggested. Weighted data is not supported for automated bin size selection.

`'auto'`

Maximum of the 'sturges' and 'fd' estimators. Provides good all around performance.

`'fd' (Freedman Diaconis Estimator)`

Robust (resilient to outliers) estimator that takes into account data variability and data size.

`'doane'`

An improved version of Sturges' estimator that works better with non-normal datasets.

`'scott'`

Less robust estimator that takes into account data variability and data size.

`'rice'`

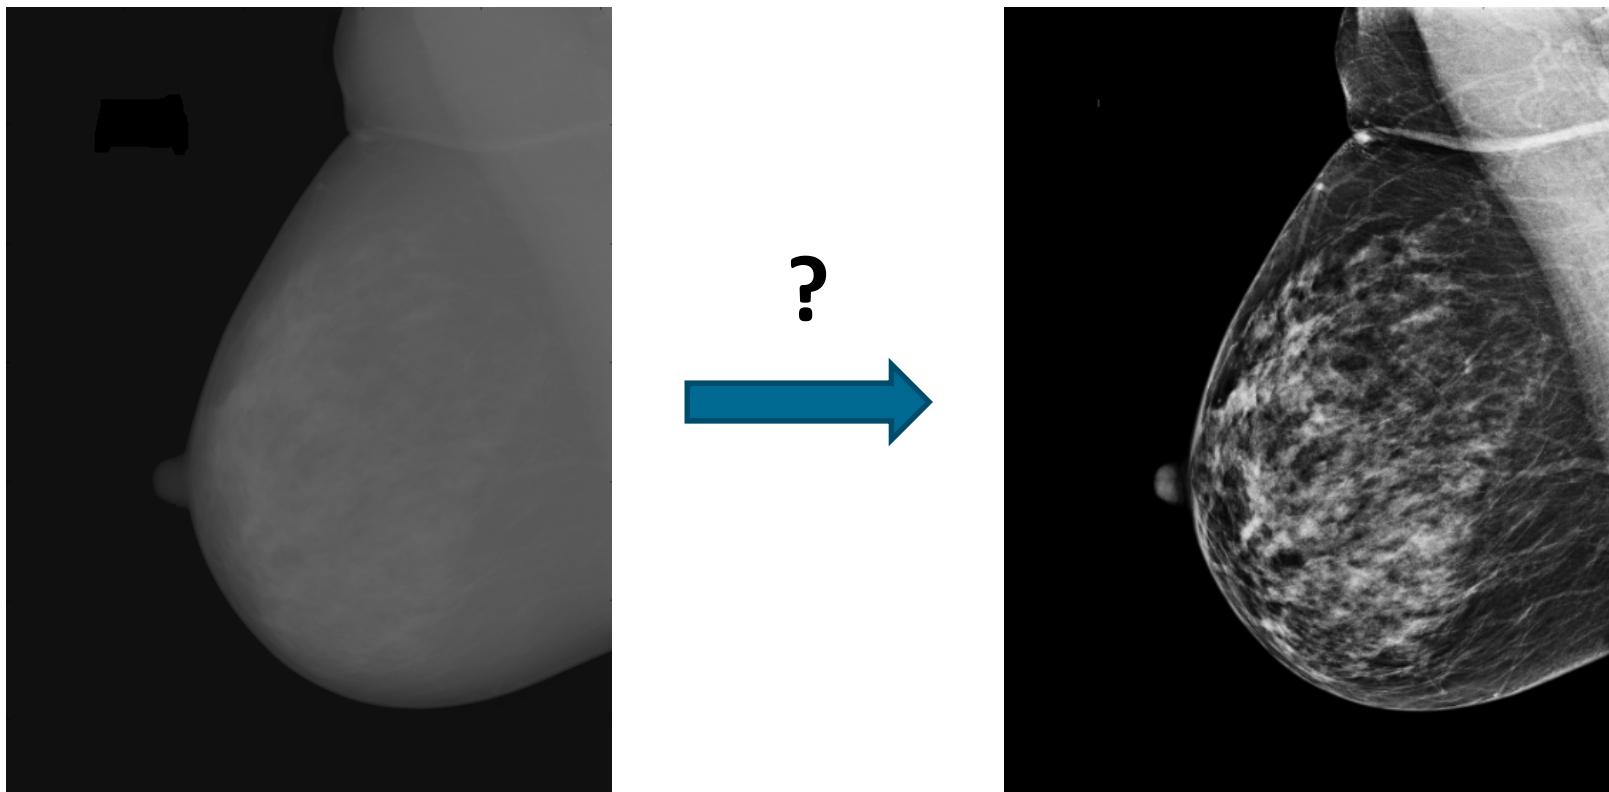
Estimator does not take variability into account, only data size. Commonly overestimates number of bins required.

`'sturges'`

R's default method, only accounts for data size. Only optimal for gaussian data and underestimates number of bins for large non-gaussian datasets.

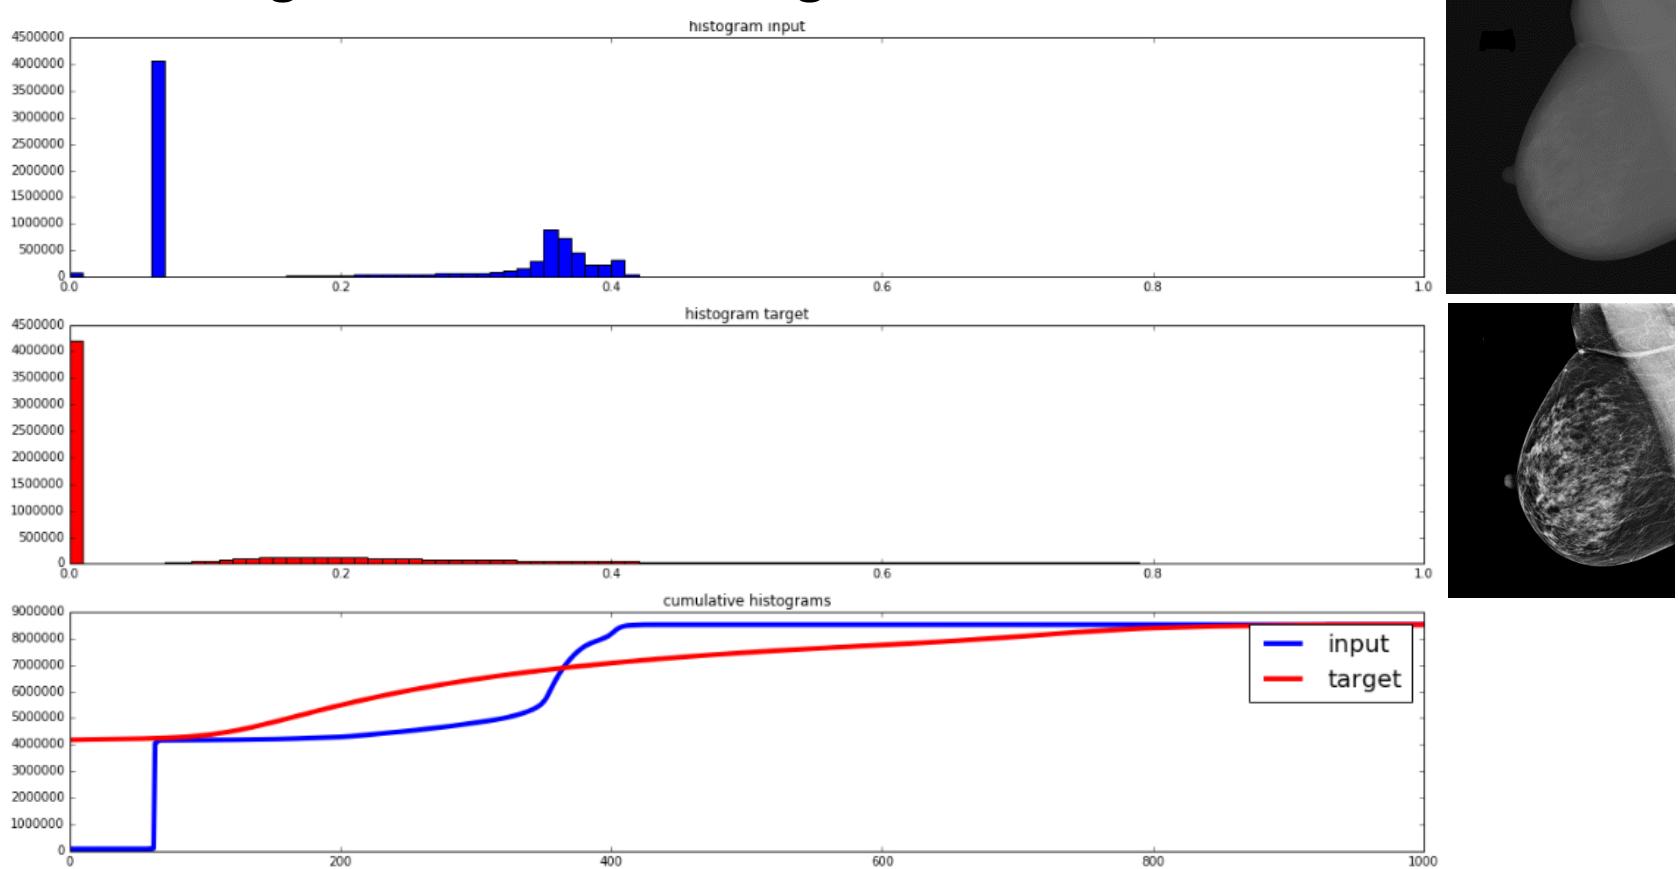
Histogram matching

- Transform the intensity of **input** image to match the histogram of a **target** image

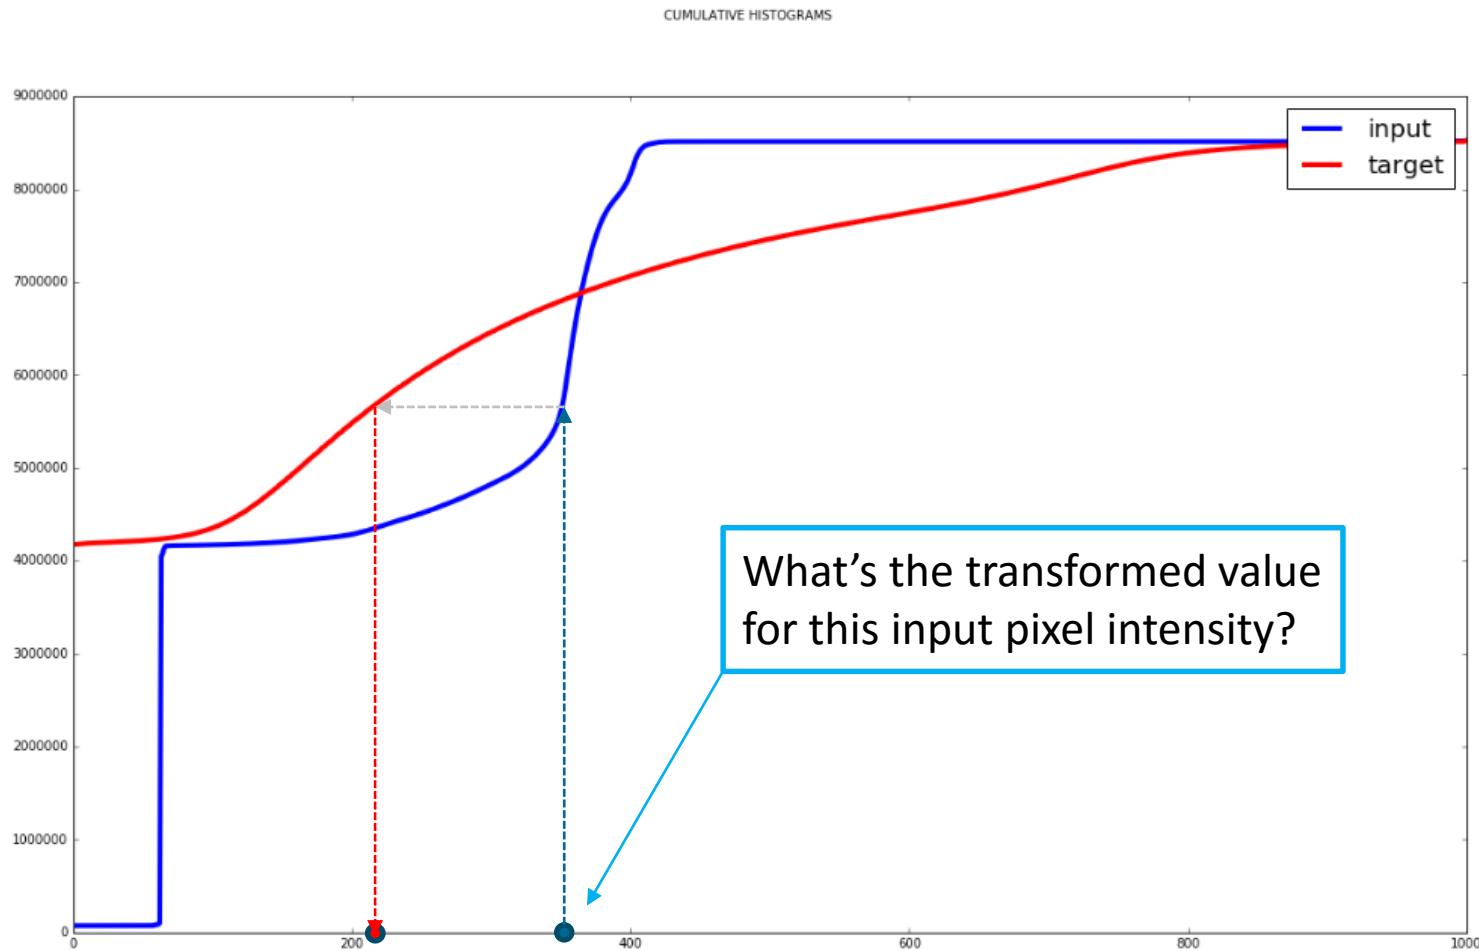


Histogram matching

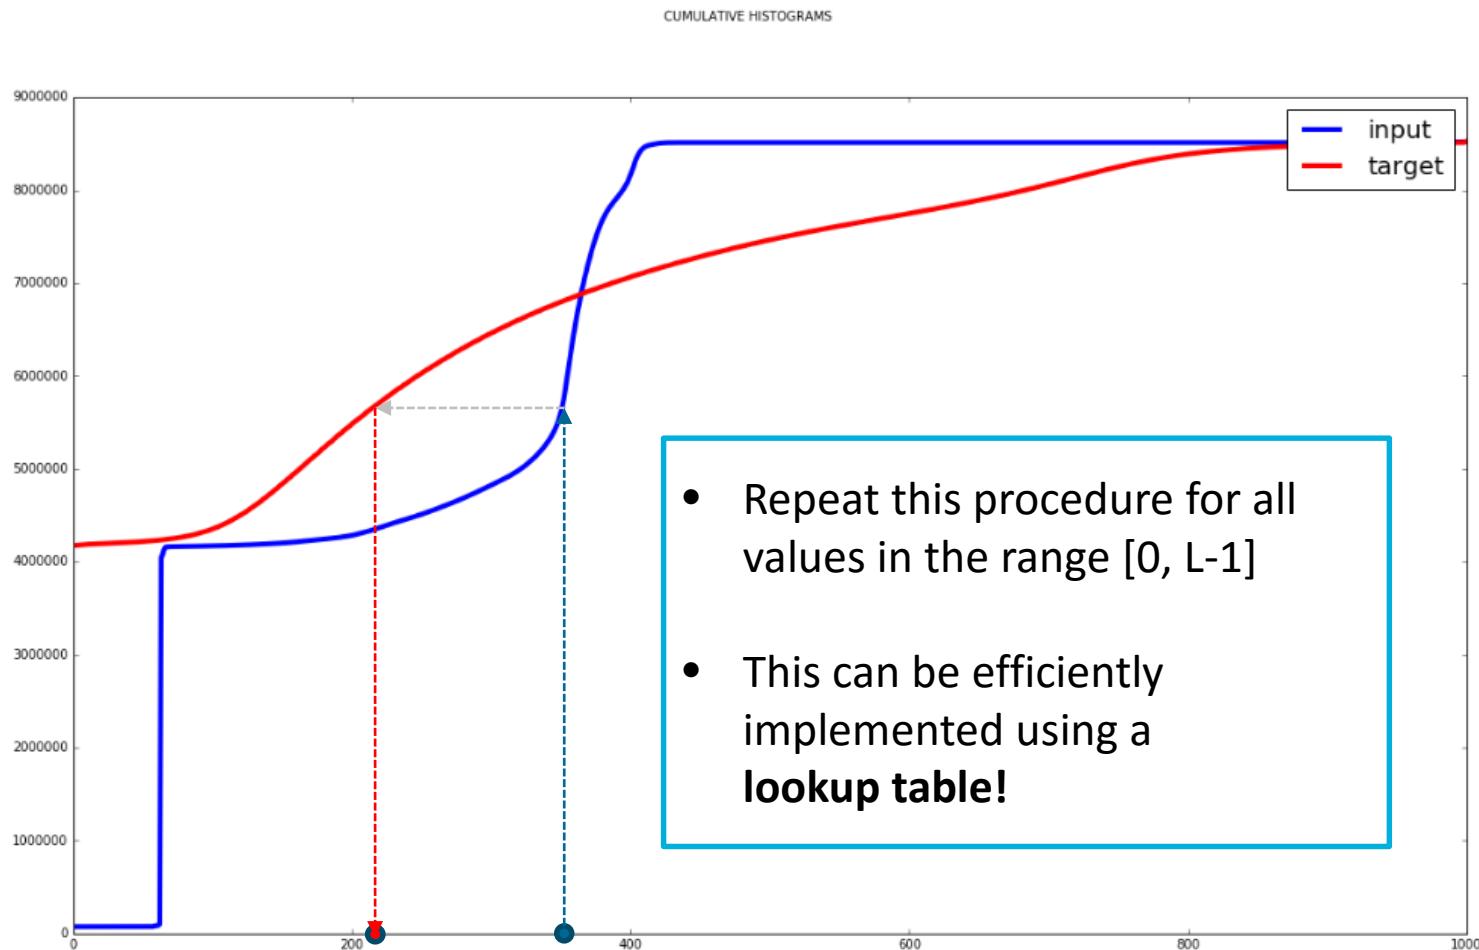
- The idea is to compute the **Cumulative Distribution Function** of the histogram of the two images



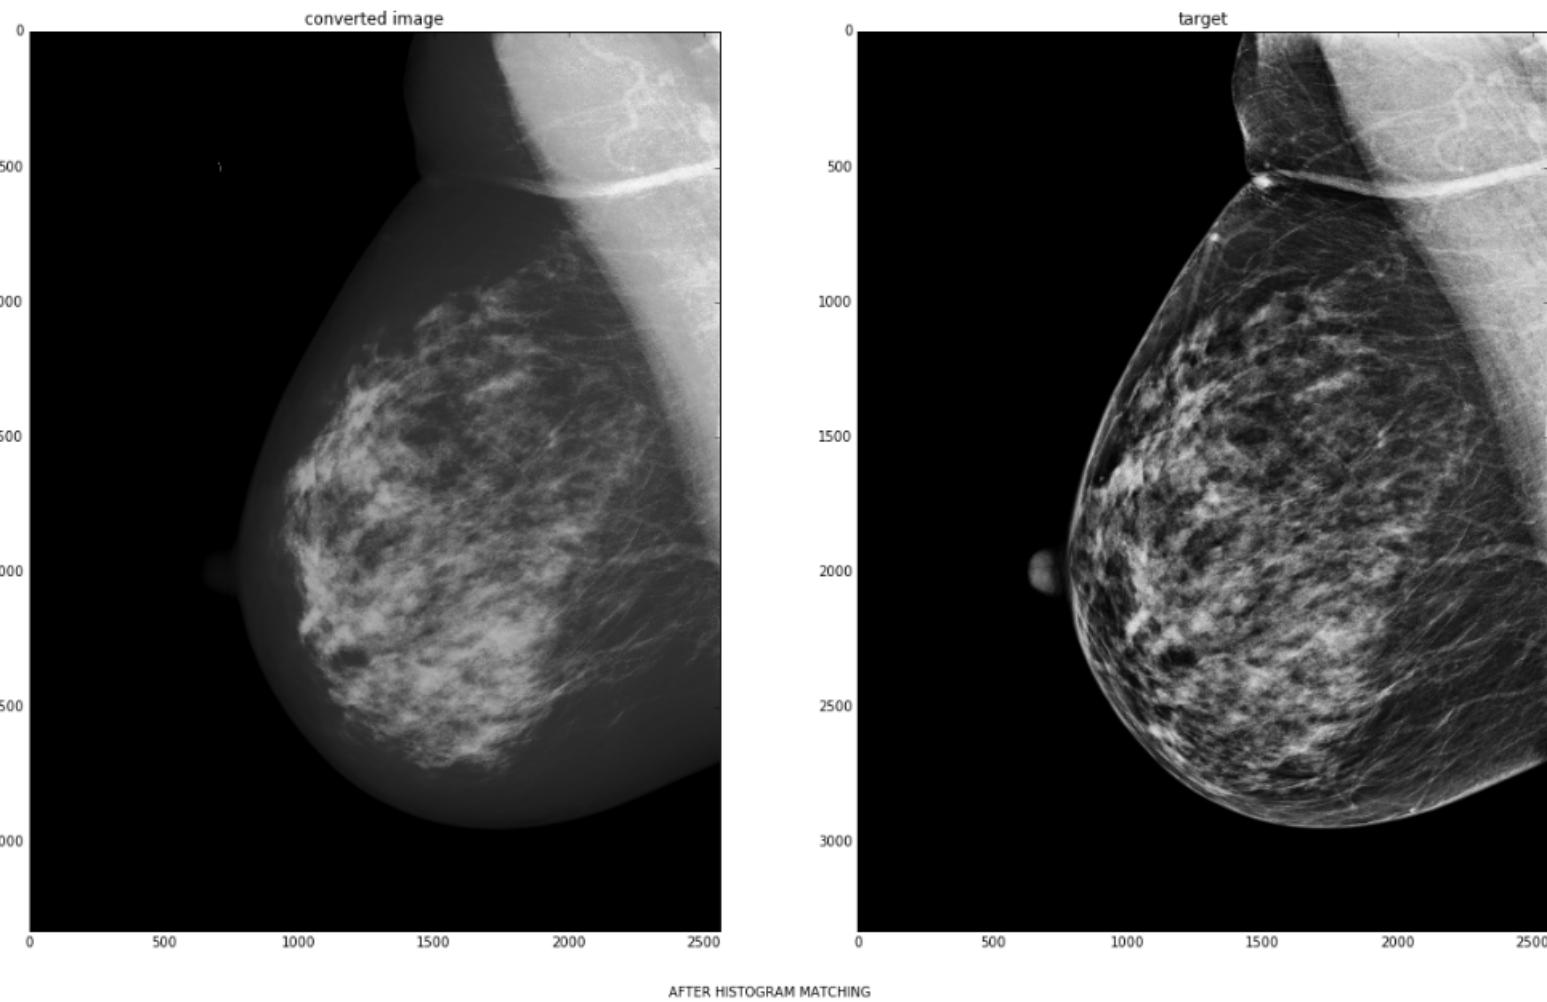
Histogram matching



Histogram matching



Histogram matching



In Python



[Scipy.org](#) [Docs](#) [NumPy v1.10 Manual](#) [NumPy Reference](#) [Routines](#) [Mathematical functions](#)

numpy.cumsum

`numpy.cumsum(a, axis=None, dtype=None, out=None)`

[\[source\]](#)

Return the cumulative sum of the elements along a given axis.

Parameters: `a : array_like`

Input array.

`axis : int, optional`

Axis along which the cumulative sum is computed. The default (None) is to compute the cumsum over the flattened array.

`dtype : dtype, optional`

Type of the returned array and of the accumulator in which the elements are summed. If `dtype` is not specified, it defaults to the `dtype` of `a`, unless `a` has an integer `dtype` with a precision less than that of the default platform integer. In that case, the default platform integer is used.

`out : ndarray, optional`

Alternative output array in which to place the result. It must have the same shape and buffer length as the expected output but the type will be cast if necessary.

See doc.ufuncs (Section "Output arguments") for more details.

Returns:

`cumsum_along_axis : ndarray`.

A new array holding the result is returned unless `out` is specified, in which case a reference to `out` is returned. The result has the same size as `a`, and the same shape as `a` if `axis` is not None or `a` is a 1-d array.

See also:

`sum` Sum array elements.

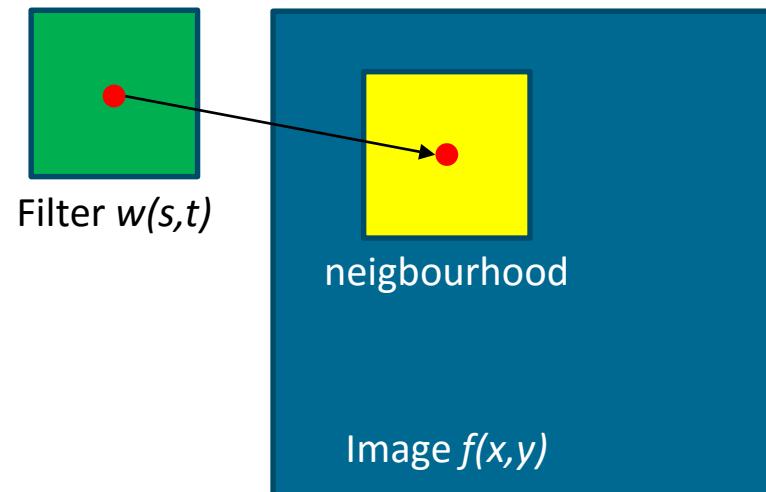
`trapz` Integration of array values using the composite trapezoidal rule.

`diff` Calculate the n-th order discrete difference along given axis.

Filtering

Spatial Filtering

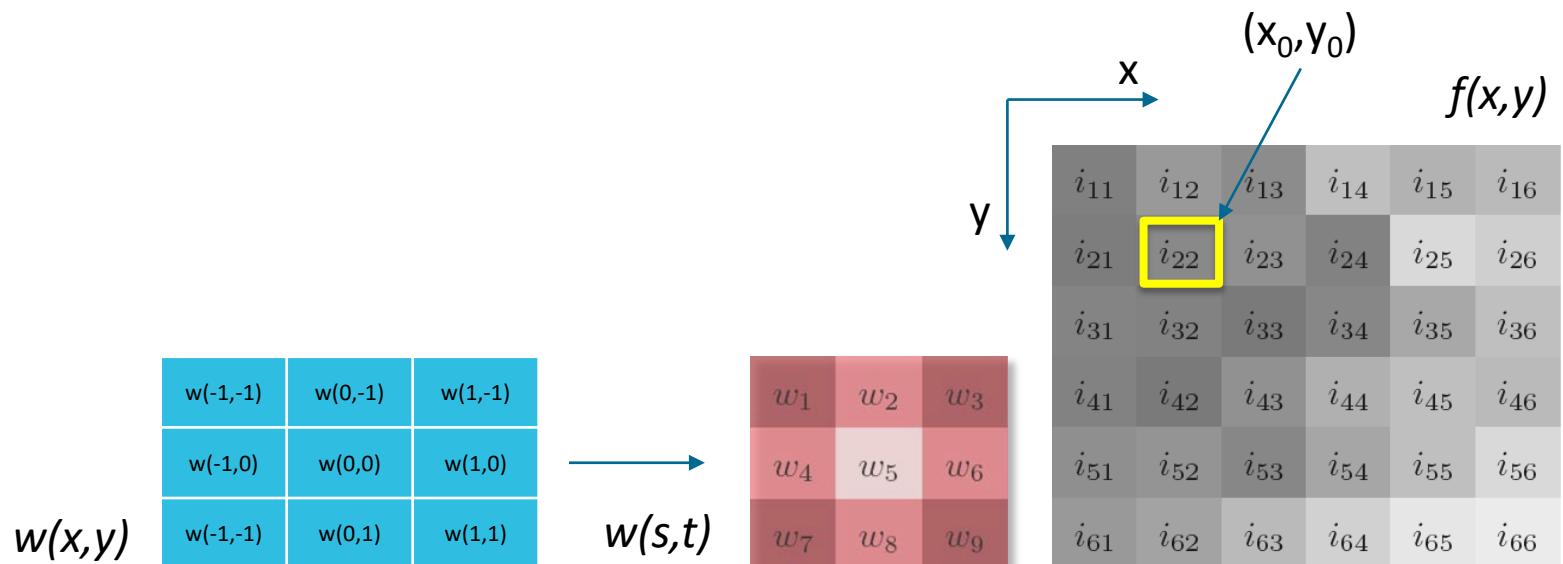
- Some operations work with the values of the image pixels in a **neighborhood**, and the corresponding values of a *subimage* that has the same dimensions as the neighborhood
- This *subimage* is called a:
 - Filter
 - Mask
 - Kernel
 - Template
 - Window
- The **values** in a *kernel* are referred to as **coefficients** rather than pixels



Correlation at one position (x_0, y_0)

- Filter operation (3x3 filter)

$$g(x_0, y_0) = w(-1, -1)f(x_0 - 1, y_0 - 1) + w(0, -1)f(x_0, y_0 - 1) + w(1, -1)f(x_0 + 1, y_0 - 1) + \\ w(-1, 0)f(x_0 - 1, y_0) + w(0, 0)f(x_0, y_0) + \boxed{w(1, 0)f(x_0 + 1, y_0)} + \\ w(-1, 1)f(x_0 - 1, y_0 + 1) + w(0, 1)f(x_0, y_0 + 1) + w(1, 1)f(x_0 + 1, y_0 + 1)$$



Correlation at each position

- While moving the filter, **at each position**

$$g(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)$$

- Multiply values of overlapping locations
- Sum all multiplication results

w_{11}	w_{12}	w_{13}	i_{14}	i_{15}	i_{16}
w_{21}	w_{22}	w_{23}	i_{24}	i_{25}	i_{26}
w_{31}	w_{32}	w_{33}	i_{34}	i_{35}	i_{36}
i_{41}	i_{42}	i_{43}	i_{44}	i_{45}	i_{46}
i_{51}	i_{52}	i_{53}	i_{54}	i_{55}	i_{56}
i_{61}	i_{62}	i_{63}	i_{64}	i_{65}	i_{66}

Intuition to Spatial Filtering

i_{11}	i_{12}	i_{13}	i_{14}	i_{15}	i_{16}
i_{21}	i_{22}	i_{23}	i_{24}	i_{25}	i_{26}
i_{31}	i_{32}	i_{33}	i_{34}	i_{35}	i_{36}
i_{41}	i_{42}	i_{43}	i_{44}	i_{45}	i_{46}
i_{51}	i_{52}	i_{53}	i_{54}	i_{55}	i_{56}
i_{61}	i_{62}	i_{63}	i_{64}	i_{65}	i_{66}

Move filter over image

w_1	w_2	w_3	i_{14}	i_{15}	i_{16}
w_4	w_5	w_6	i_{24}	i_{25}	i_{26}
w_7	w_8	w_9	i_{34}	i_{35}	i_{36}
i_{41}	i_{42}	i_{43}	i_{44}	i_{45}	i_{46}
i_{51}	i_{52}	i_{53}	i_{54}	i_{55}	i_{56}
i_{61}	i_{62}	i_{63}	i_{64}	i_{65}	i_{66}

i_{11}	w_1	w_2	w_3	i_{15}	i_{16}
i_{21}	w_4	w_5	w_6	i_{25}	i_{26}
i_{31}	w_7	w_8	w_9	i_{35}	i_{36}
i_{41}	i_{42}	i_{43}	i_{44}	i_{45}	i_{46}
i_{51}	i_{52}	i_{53}	i_{54}	i_{55}	i_{56}
i_{61}	i_{62}	i_{63}	i_{64}	i_{65}	i_{66}

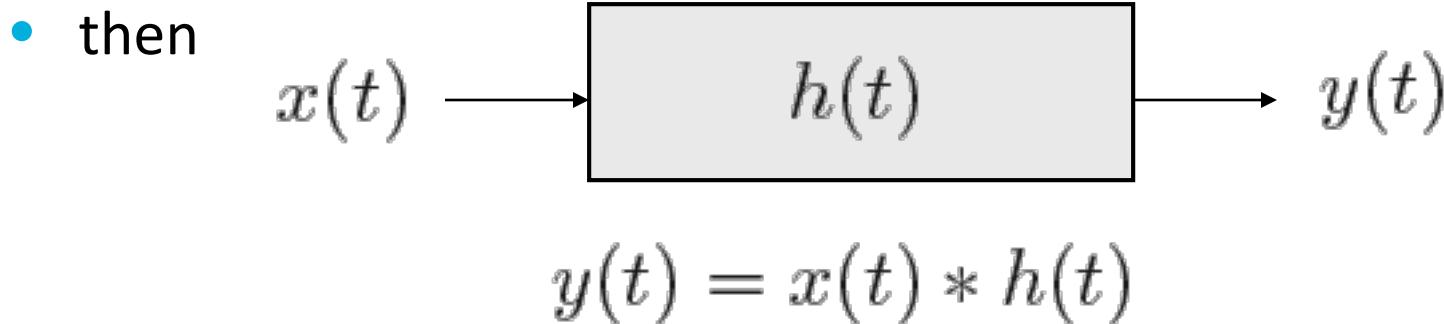
i_{11}	i_{12}	w_1	w_2	w_3	i_{16}
i_{21}	i_{22}	w_4	w_5	w_6	i_{26}
i_{31}	i_{32}	w_7	w_8	w_9	i_{36}
i_{41}	i_{42}	i_{43}	i_{44}	i_{45}	i_{46}
i_{51}	i_{52}	i_{53}	i_{54}	i_{55}	i_{56}
i_{61}	i_{62}	i_{63}	i_{64}	i_{65}	i_{66}

i_{11}	i_{12}	i_{13}	w_1	w_2	w_3
i_{21}	i_{22}	i_{23}	w_4	w_5	w_6
i_{31}	i_{32}	i_{33}	w_7	w_8	w_9
i_{41}	i_{42}	i_{43}	i_{44}	i_{45}	i_{46}
i_{51}	i_{52}	i_{53}	i_{54}	i_{55}	i_{56}
i_{61}	i_{62}	i_{63}	i_{64}	i_{65}	i_{66}

i_{11}	i_{12}	i_{13}	i_{14}	i_{15}	i_{16}
w_1	w_2	w_3	i_{24}	i_{25}	i_{26}
w_4	w_5	w_6	i_{34}	i_{35}	i_{36}
w_7	w_8	w_9	i_{44}	i_{45}	i_{46}
i_{51}	i_{52}	i_{53}	i_{54}	i_{55}	i_{56}
i_{61}	i_{62}	i_{63}	i_{64}	i_{65}	i_{66}

Convolution

- Imagine a system with
 - input signal $x(t)$
 - transfer function $h(t)$
 - output signal $y(t)$



Correlation vs. Convolution

- Discrete **Convolution 2D**

$$g[i, j] = \sum_{s=-a}^a \sum_{t=-b}^b w[s, t] f[i - s, j - t]$$

$$y(t) = \int_{-\infty}^{\infty} x(\tau) h(t - \tau) d\tau$$

- Discrete **Correlation 2D**

$$g[i, j] = \sum_{s=-a}^a \sum_{t=-b}^b w[s, t] f[i + s, j + t]$$

- Equivalent to first rotating the filter 180 degrees and using correlation
- Correlation and convolution give the same result if the filter used is symmetric!

Correlation vs. Convolution

w_9	w_8	w_7	i_{14}	i_{15}	i_{16}
w_6	w_5	w_4	i_{24}	i_{25}	i_{26}
w_3	w_2	w_1	i_{34}	i_{35}	i_{36}
i_{41}	i_{42}	i_{43}	i_{44}	i_{45}	i_{46}
i_{51}	i_{52}	i_{53}	i_{54}	i_{55}	i_{56}
i_{61}	i_{62}	i_{63}	i_{64}	i_{65}	i_{66}

In Python

 SciPy.org 

[Scipy.org](#) [Docs](#) [SciPy v0.18.1 Reference Guide](#) [Signal processing \(`scipy.signal`\)](#)

scipy.signal.convolve2d

`scipy.signal.convolve2d(in1, in2, mode='full', boundary='fill', fillvalue=0)` [\[source\]](#)

Convolve two 2-dimensional arrays.

Convolve *in1* and *in2* with output size determined by *mode*, and boundary conditions determined by *boundary* and *fillvalue*.

Parameters:

- in1** : *array_like*
First input.
- in2** : *array_like*
Second input. Should have the same number of dimensions as *in1*. If operating in 'valid' mode, either *in1* or *in2* must be at least as large as the other in every dimension.
- mode** : *str* {'full', 'valid', 'same'}, *optional*
A string indicating the size of the output:
 - full**
The output is the full discrete linear convolution of the inputs. (Default)
 - valid**
The output consists only of those elements that do not rely on the zero-padding.
 - same**
The output is the same size as *in1*, centered with respect to the 'full' output.
- boundary** : *str* {'fill', 'wrap', 'symm'}, *optional*
A flag indicating how to handle boundaries:
 - fill**
pad input arrays with *fillvalue*. (default)
 - wrap**
circular boundary conditions.

Filters

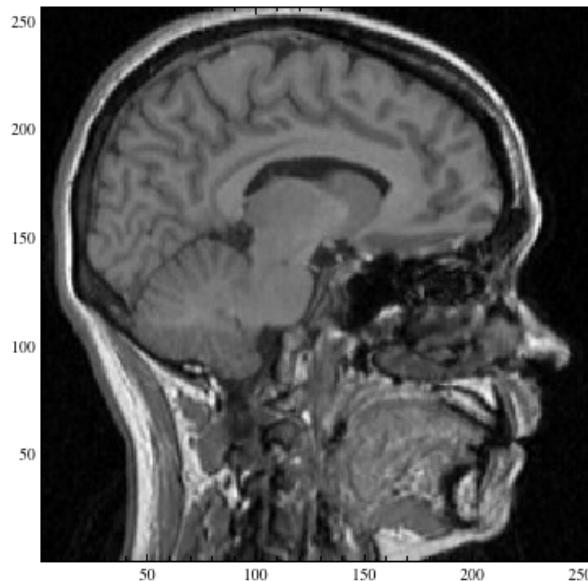
- What kind of filters can we design?
 - Image smoothing
 - Edge detection
 - Image derivative
 - Multi-scale analysis
 - Blob detection

Image smoothing

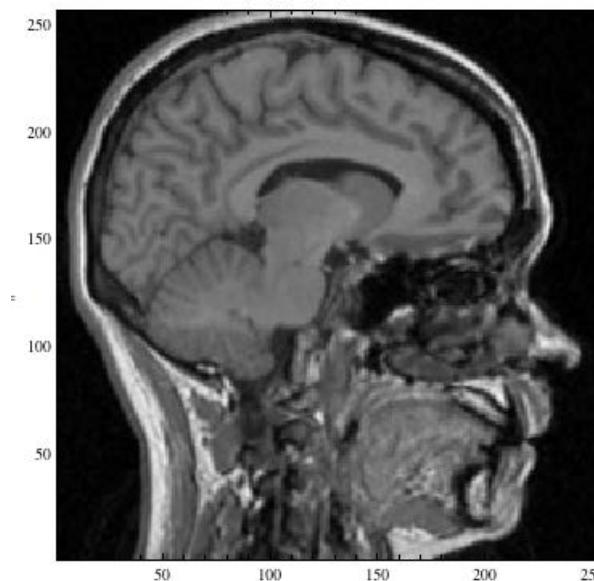
- Smoothing filters are used for **blurring** and for **noise reduction**
- Based on **averaging** of brightness values in some neighborhood
- Average intensities – result is blurred image, less details

$$\frac{1}{9} \begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix} \quad \frac{1}{25} \begin{pmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{pmatrix} \dots \text{NxN filter}$$

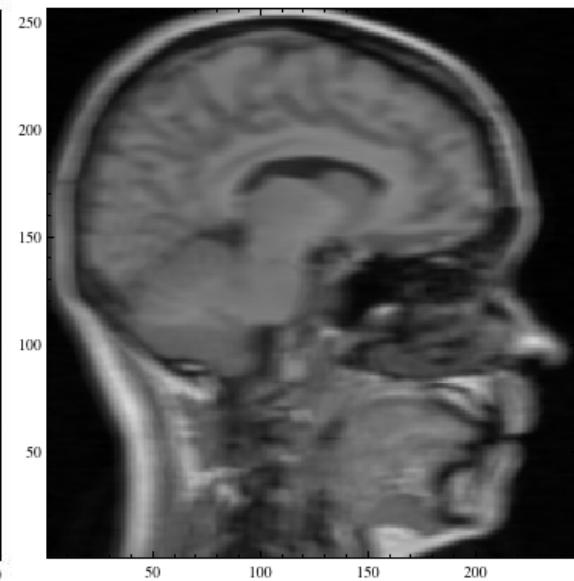
Smoothing filters



Original

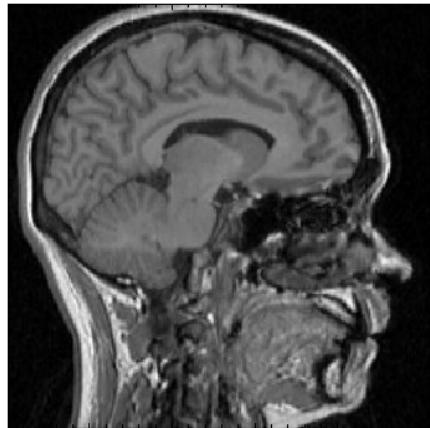


3x3 filter

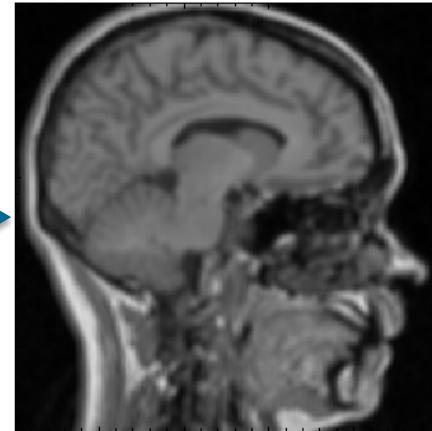


9x9 filter

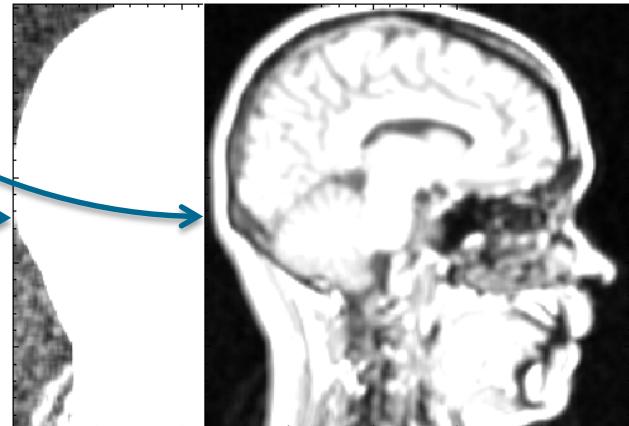
Effect of normalized smoothing kernel



normalized



non-normalized



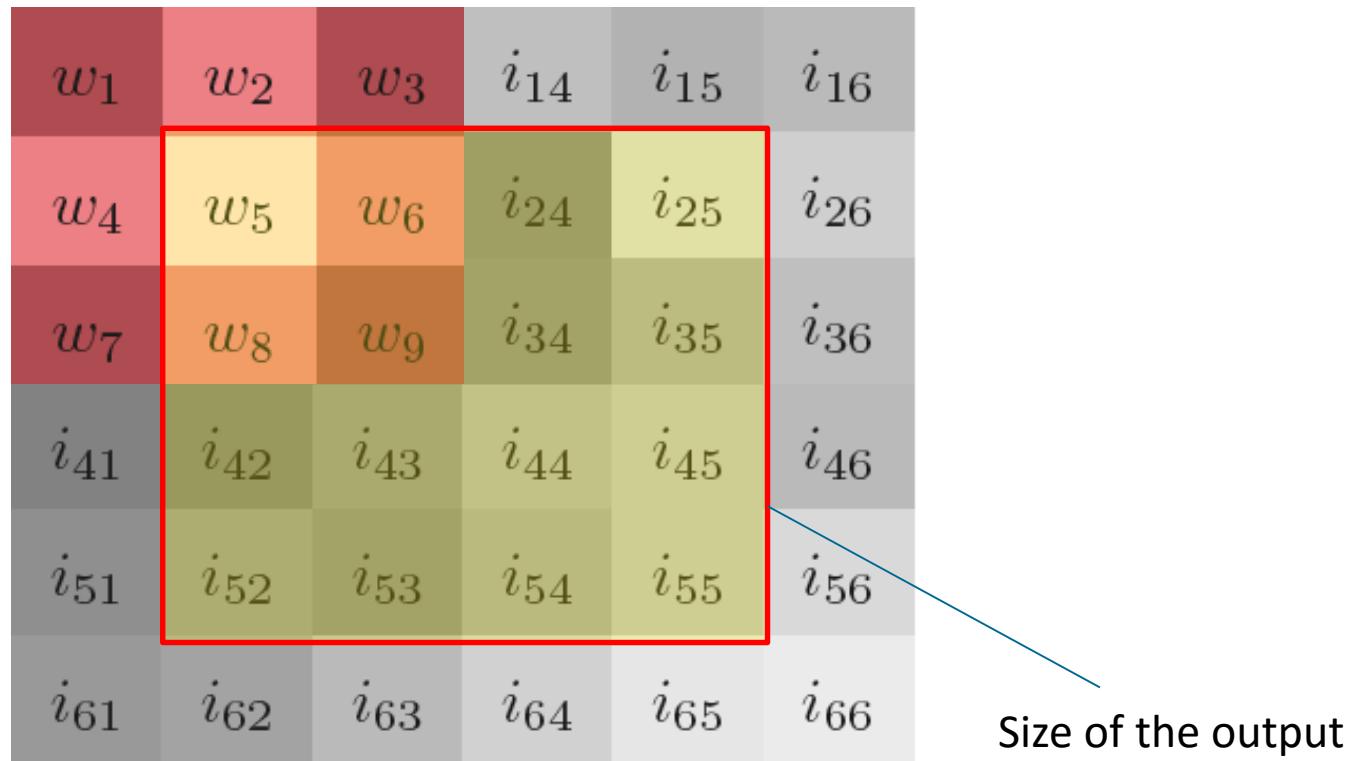
Borders

- Do you see any problems at image borders?

- Try position (0,0)

w_1	w_2	w_3				
w_4	w_5	w_6	i_{13}	i_{14}	i_{15}	i_{16}
w_7	w_8	w_9	i_{23}	i_{24}	i_{25}	i_{26}
i_{31}	i_{32}	i_{33}	i_{34}	i_{35}	i_{36}	
i_{41}	i_{42}	i_{43}	i_{44}	i_{45}	i_{46}	
i_{51}	i_{52}	i_{53}	i_{54}	i_{55}	i_{56}	
i_{61}	i_{62}	i_{63}	i_{64}	i_{65}	i_{66}	

Valid convolution



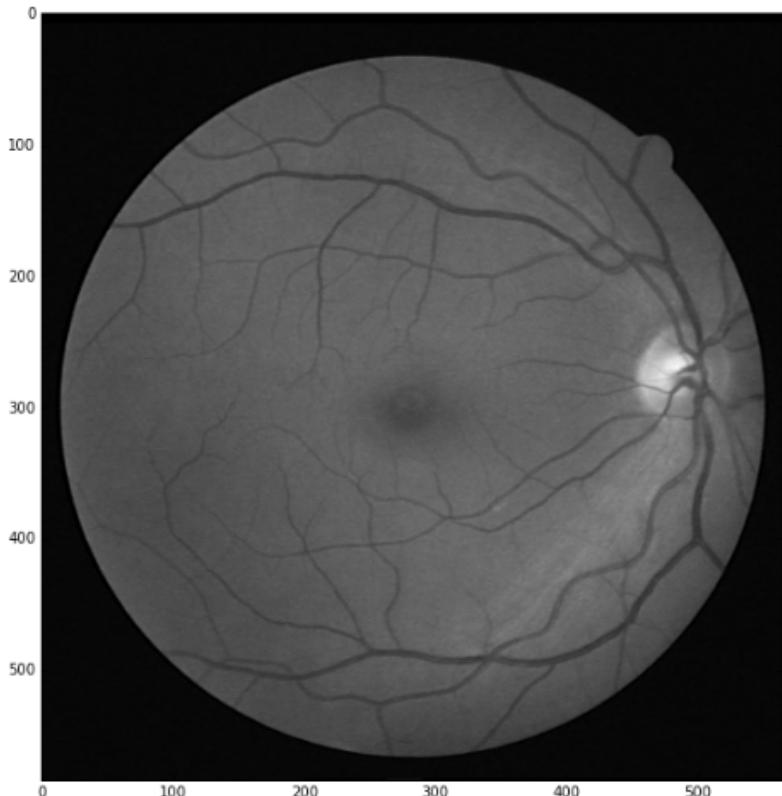
Zero padding

w_1	w_2	w_3	0	0	0	0	0
w_4	w_5	w_6	i_{13}	i_{14}	i_{15}	i_{16}	0
w_7	w_8	w_9	i_{23}	i_{24}	i_{25}	i_{26}	0
0	i_{31}	i_{32}	i_{33}	i_{34}	i_{35}	i_{36}	0
0	i_{41}	i_{42}	i_{43}	i_{44}	i_{45}	i_{46}	0
0	i_{51}	i_{52}	i_{53}	i_{54}	i_{55}	i_{56}	0
0	i_{61}	i_{62}	i_{63}	i_{64}	i_{65}	i_{66}	0
0	0	0	0	0	0	0	0

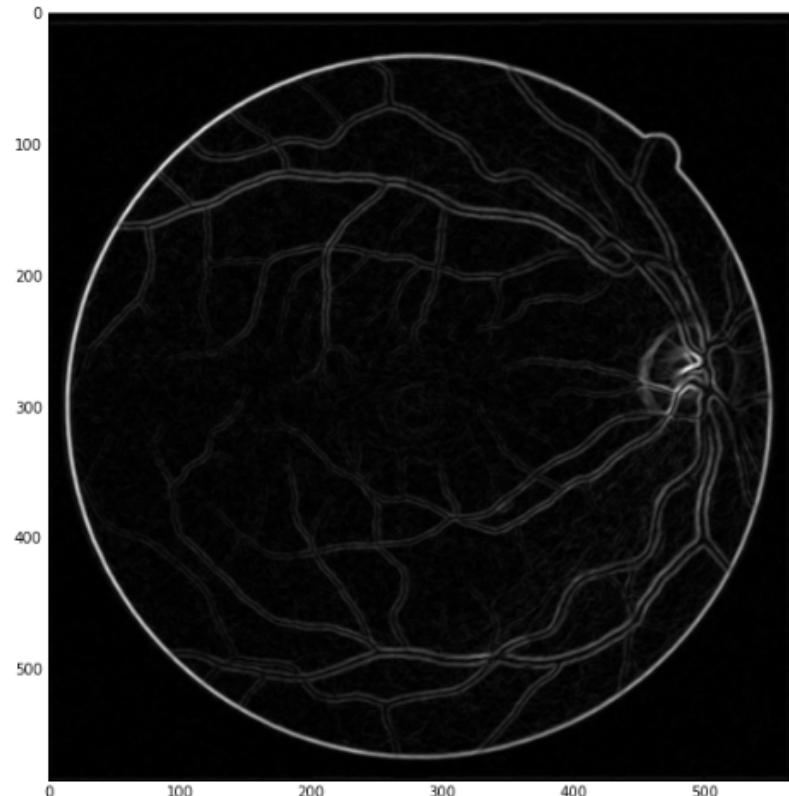
Cyclic padding

w_1	w_2	w_3	i_{63}	i_{64}	i_{65}	i_{66}	i_{61}
w_4	w_5	w_6	i_{13}	i_{14}	i_{15}	i_{16}	i_{11}
w_7	w_8	w_9	i_{23}	i_{24}	i_{25}	i_{26}	i_{21}
i_{36}	i_{31}	i_{32}	i_{33}	i_{34}	i_{35}	i_{36}	i_{31}
i_{46}	i_{41}	i_{42}	i_{43}	i_{44}	i_{45}	i_{46}	i_{41}
i_{56}	i_{51}	i_{52}	i_{53}	i_{54}	i_{55}	i_{56}	i_{51}
i_{66}	i_{61}	i_{62}	i_{63}	i_{64}	i_{65}	i_{66}	i_{61}
i_{16}	i_{11}	i_{12}	i_{13}	i_{14}	i_{15}	i_{16}	i_{11}

Edge detection



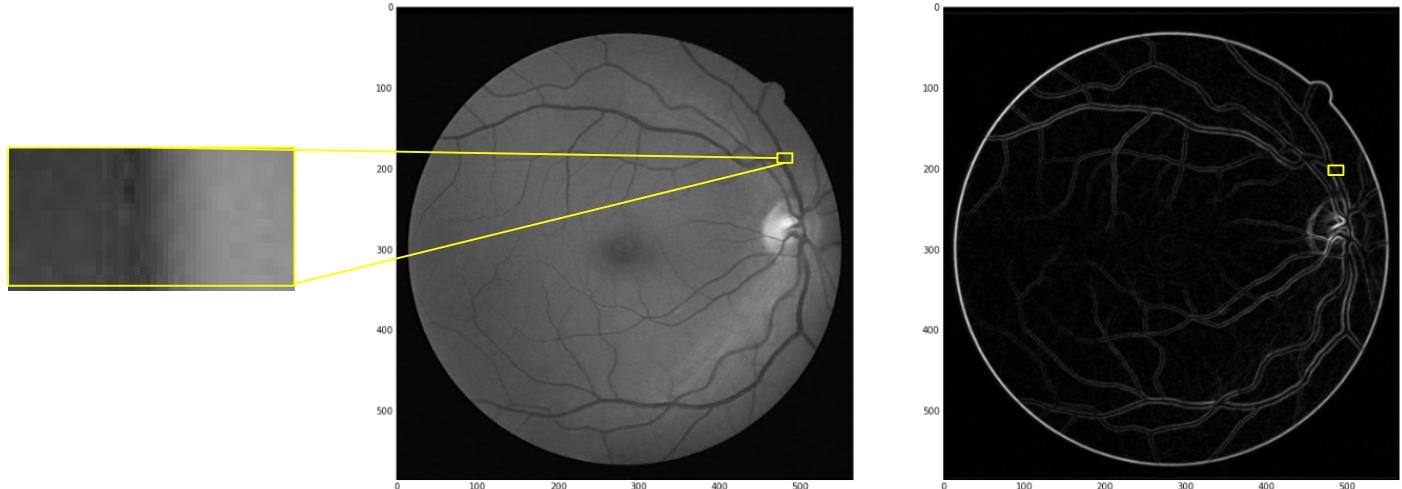
Green channels of retina fundus



How to detect edges?

- The kernel should have a sign change (-1 and +1)
- Now you measure the difference between pixels
- If the kernel looks like an edge, it responds on edges

$$\begin{matrix} -1 & \\ & 1 \end{matrix}$$



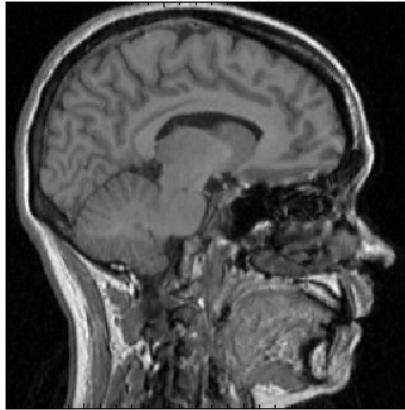
- **This is true in general:** a kernel with sign changes responds best to similar structures

Edge Detection

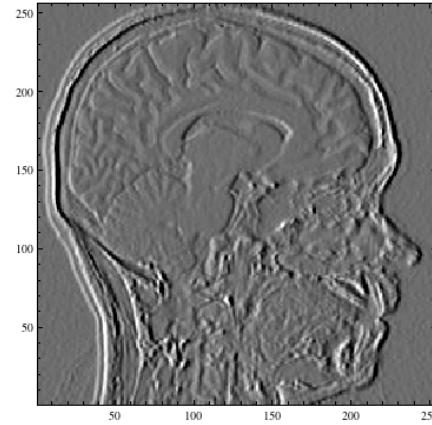
- Basic derivative filters

$$[h_x] = [h_y]^T = \begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$$

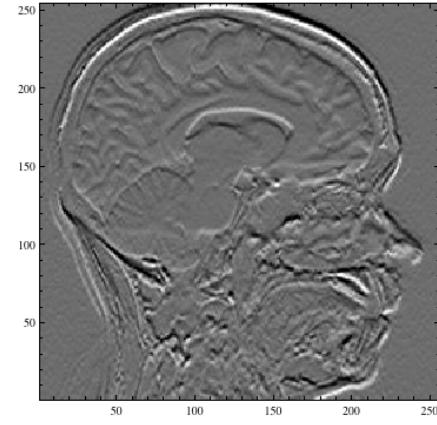
- Measure change of intensity in x or y direction



$I(x, y)$



$I(x, y) * h_x$



$I(x, y) * h_y$

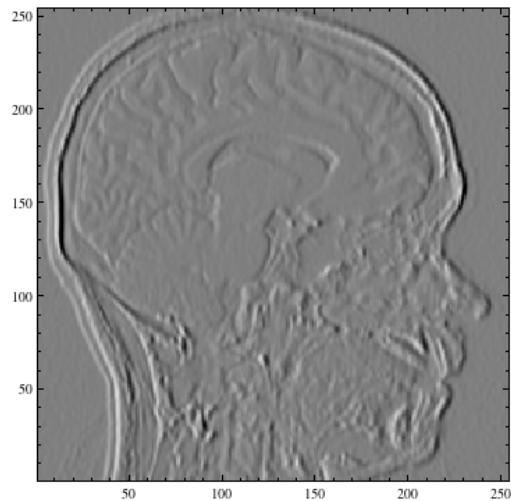
Prewitt gradient kernel

$$[h_x] = \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix} \quad [h_y] = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

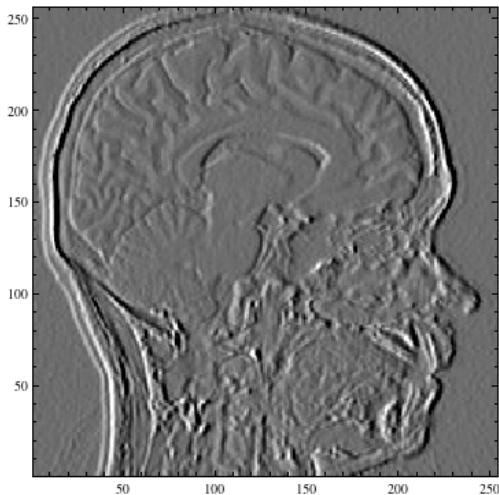
- **Derivative** in one direction, **smoothing** in the perpendicular direction
- The coefficients **sum to zero**, i.e., **zero response** in area of **constant** intensity

Example

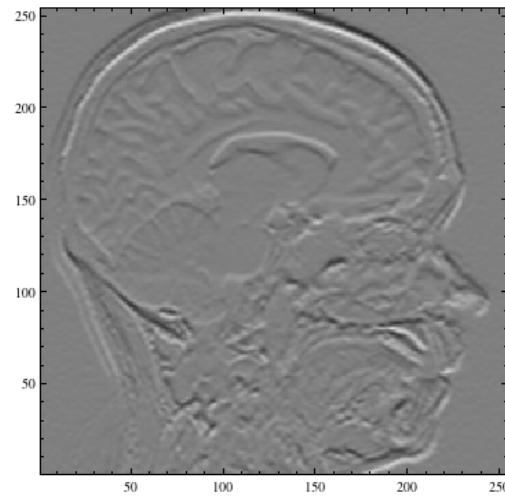
$[h_x]$



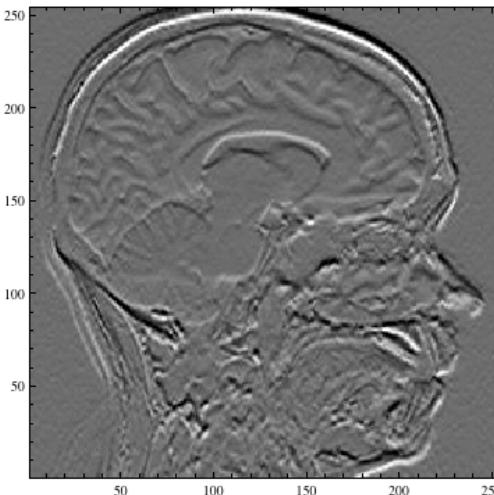
$[h_x]$



$[h_y]$ Prewitt



$[h_y]$ Basic
derivative



Sobel kernel

$$[h_x] = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

$$[h_y] = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Sobel kernel

$$[h_x] = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

$$[h_y] = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Multi-scale image analysis

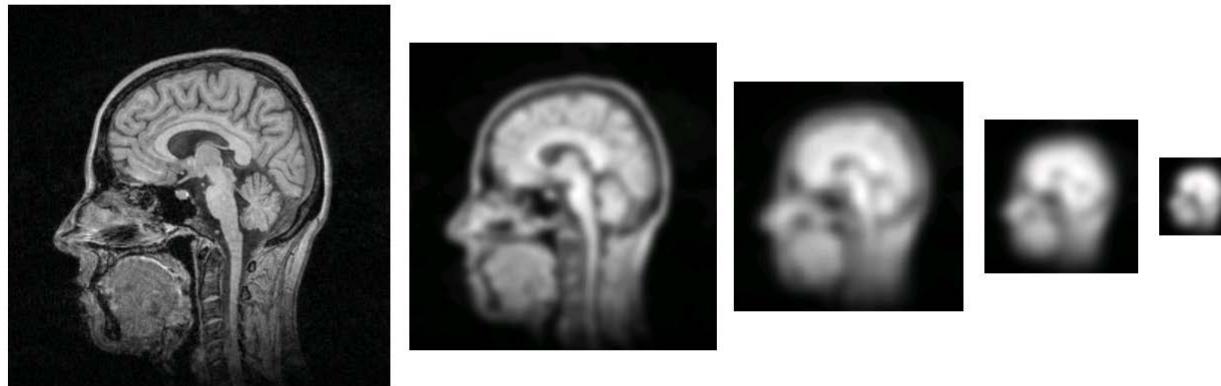
Scale space theory

- An image is a measurement of the outside world
- Images are composed of different structure at different scales

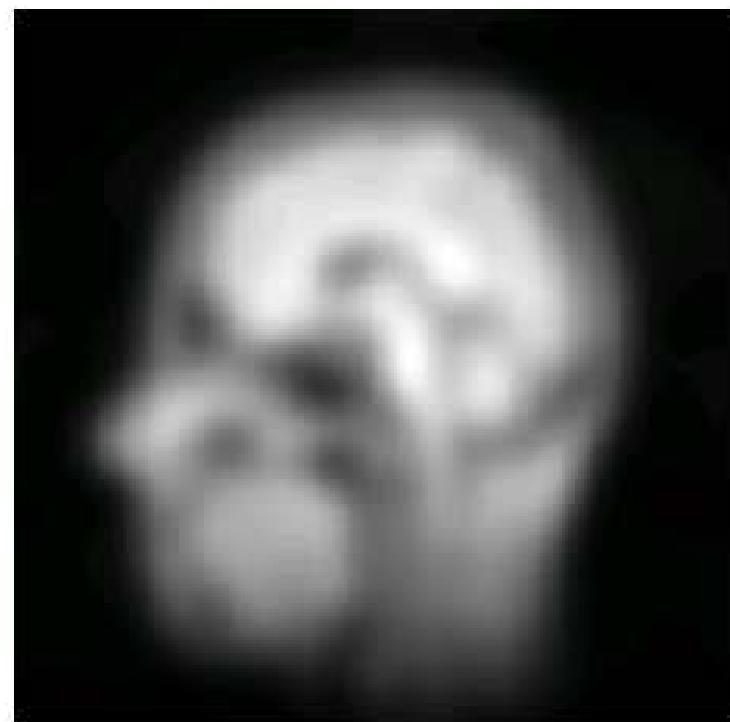
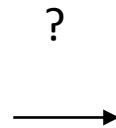
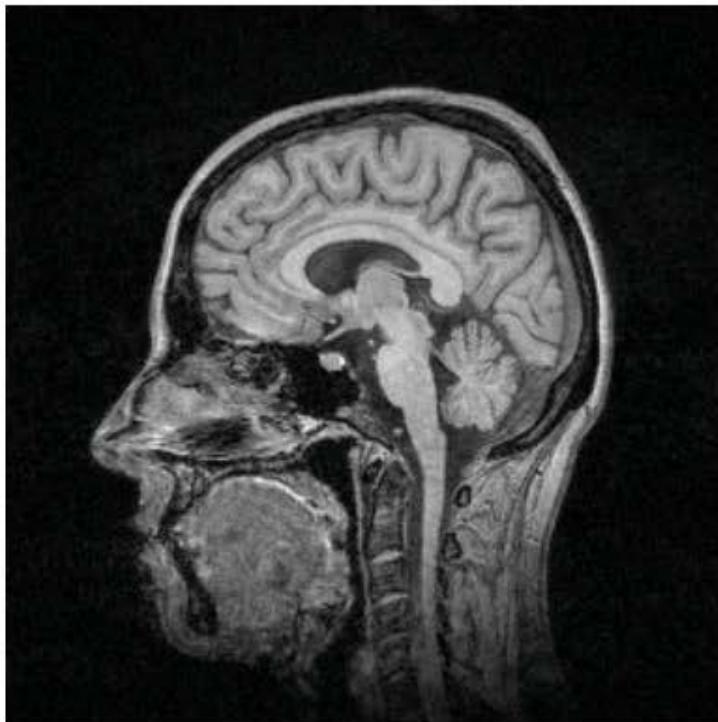


Scale space theory

- An image is a measurement of the outside world
- Images are composed of different structure at different scales
- Scale (resolution) determines what is visible in the image
- We should consider the image at all (relevant) scales

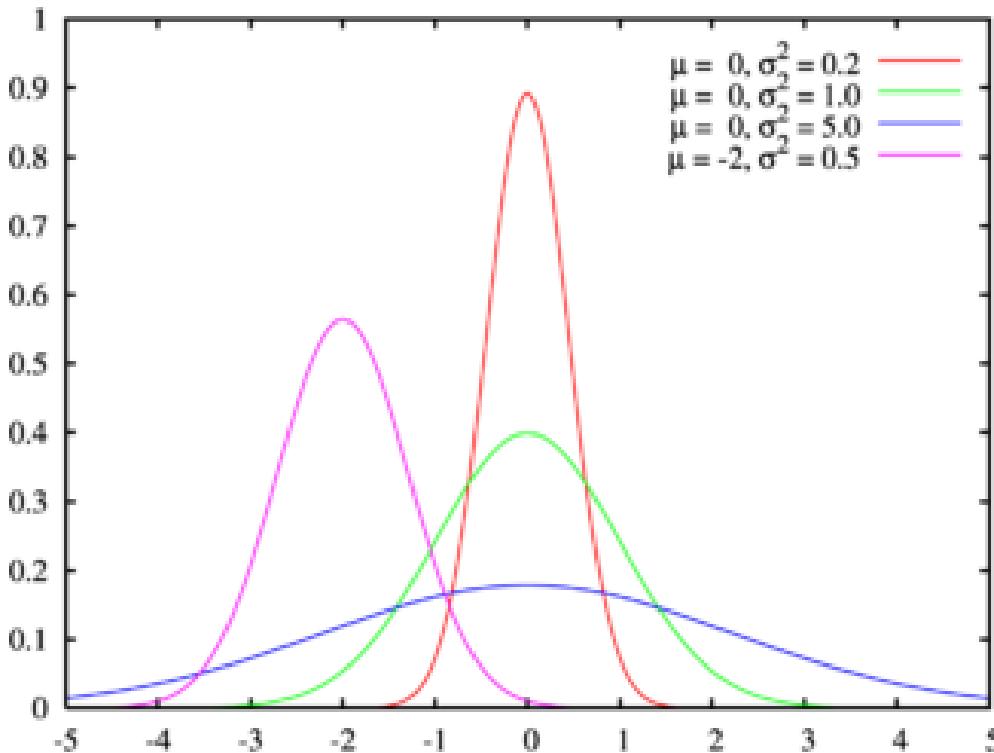


How to change scale?



Gaussian operator in 1D

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-(x-\mu)^2/(2\sigma^2)}$$



$\mu = \text{mean value}$

Controls the *position* of the “bell”

$\sigma = \text{standard deviation}$

Controls the *width* of the “bell”

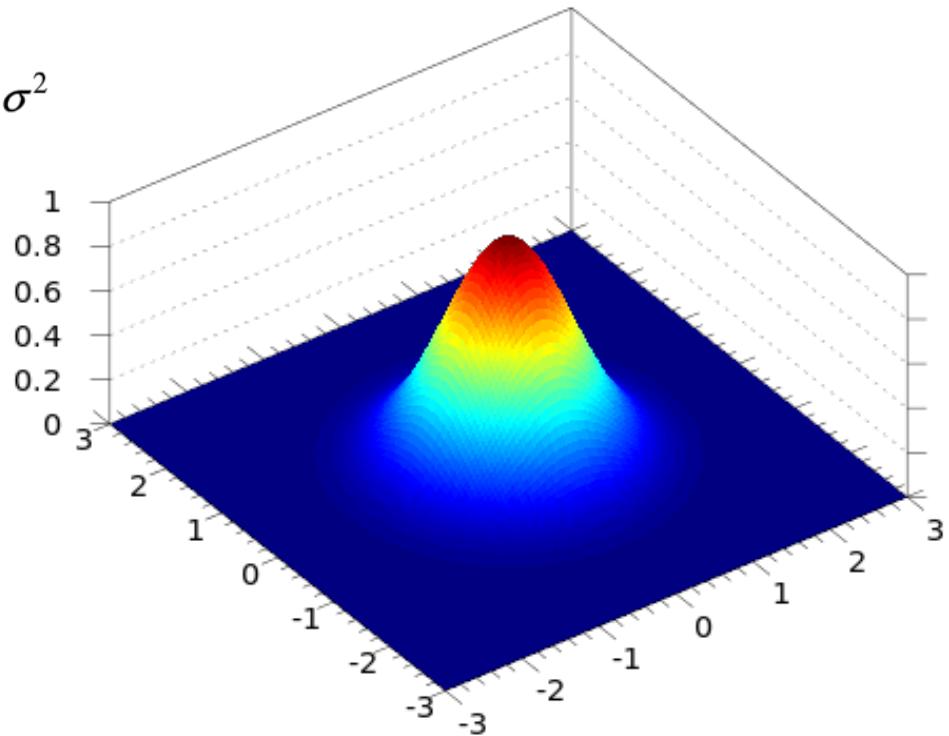
Gaussian operator in 2D

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/2\sigma^2}$$

We assume:

$$\mu_x = \mu_y = 0$$

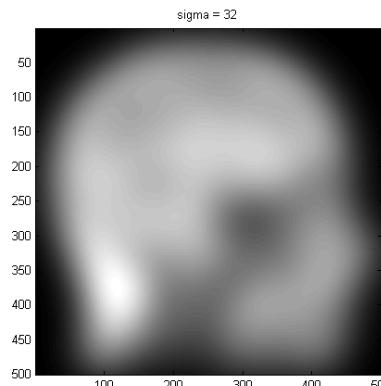
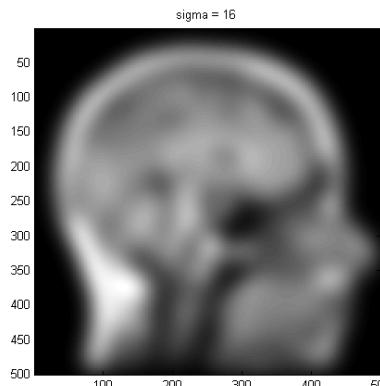
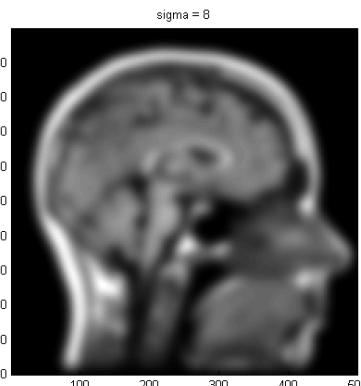
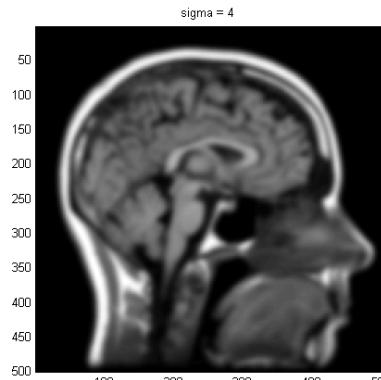
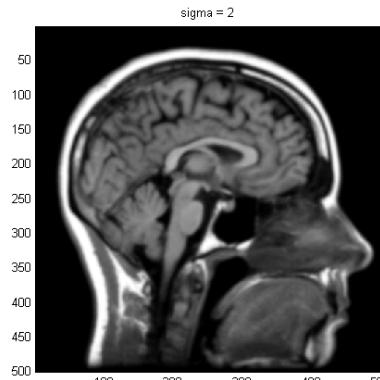
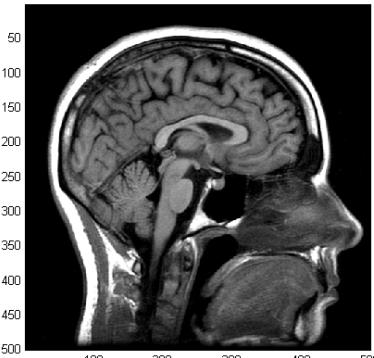
$$\sigma_x = \sigma_y = \sigma$$



- σ is **the only parameter** of the Gaussian filter
- Convolving this function with an image, **blurs** the image

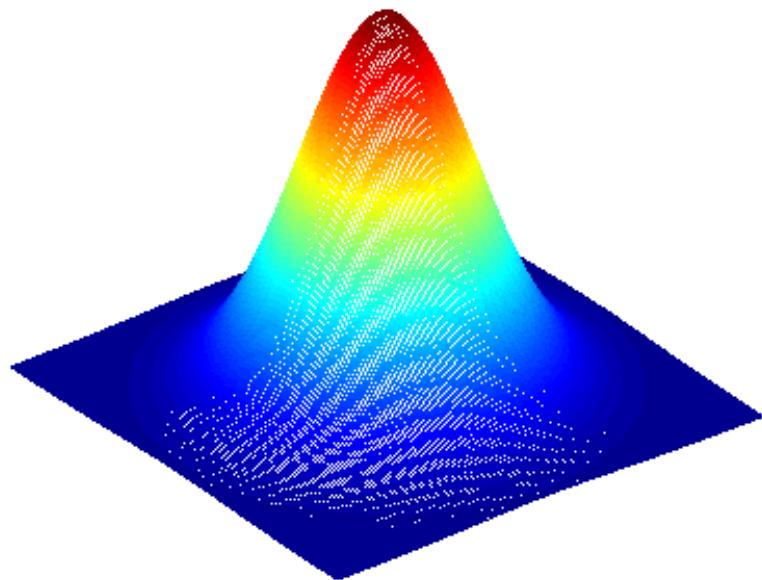
Scale space

- By changing σ and convolving with the Gaussian kernel, the image can be observed at any scale

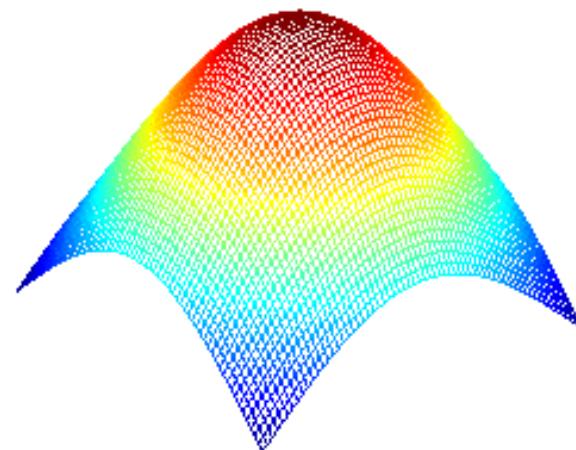


Design of a Gaussian kernel

- Sometimes you have to define the **kernel size** for a Gaussian function
- **Rule of thumb:** the Gaussian function goes to zero (approx.) after 3σ



kernel size = $[6\sigma + 1, 6\sigma + 1]$

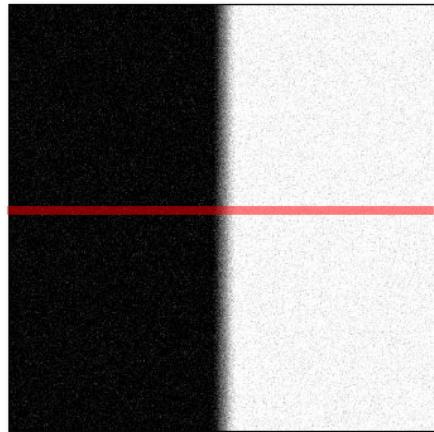


kernel size = $[2\sigma, 2\sigma]$

Derivatives and edge detection

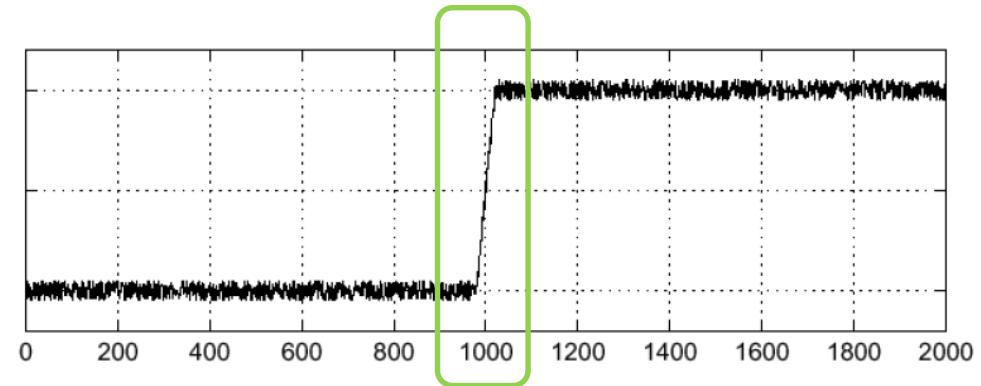
Image Derivatives

We want to detect this **edge**

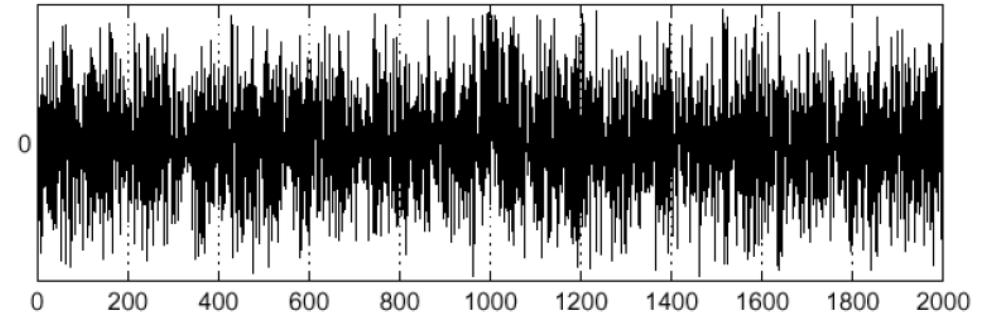


Noisy input image

$$f(x)$$



$$\frac{d}{dx}f(x)$$



We don't get much information from this

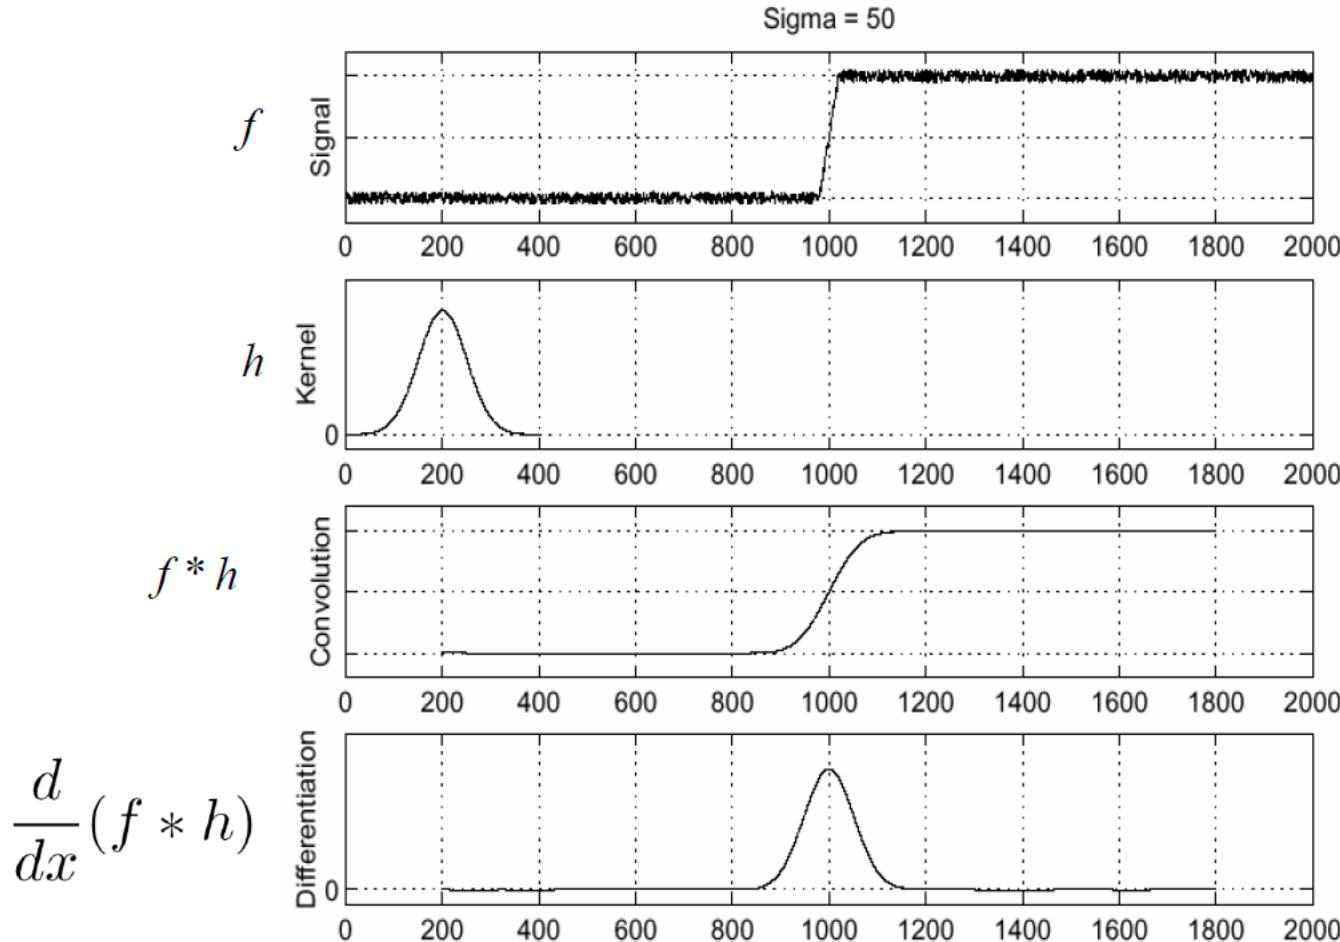
Image Derivatives using Gaussian Convolution

- **Solution:** *linearity of convolution* theorem and *scale space*
 - Convolution of $G(x)$ with derivative of $I(x)$ is the same as convolution of $I(x)$ with the derivative of $G(x)$ (*associative law*):

$$\frac{\partial I(x)}{\partial x} * G(x) = I(x) * \frac{\partial G(x)}{\partial x}$$

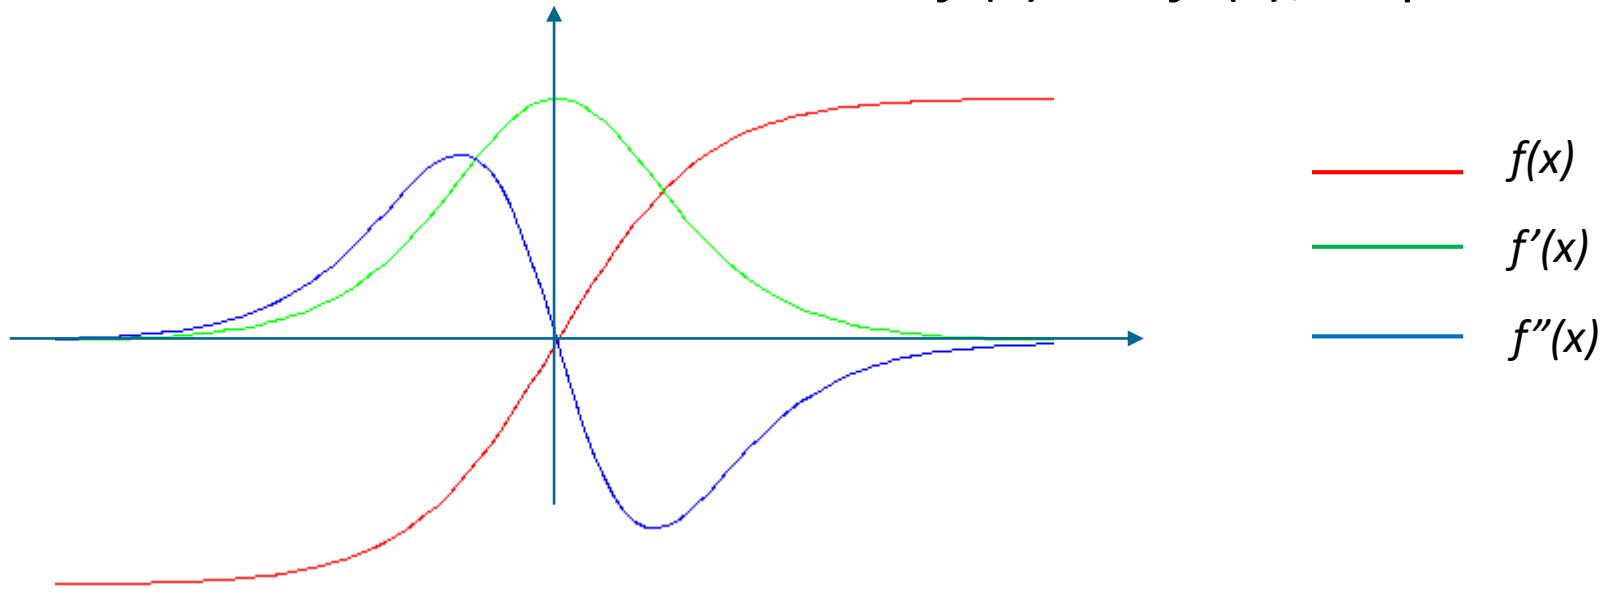
- So to compute a derivative of an image, convolve it with the derivative of a Gaussian
- Introduces a **scale dependency** (σ) in the computation of derivatives

Image Derivatives using Gaussian Convolution



Derivatives and edges

- Given an **edge function** $f(x)$, we can easily compute the first order and second order derivative $f'(x)$ and $f''(x)$, respectively

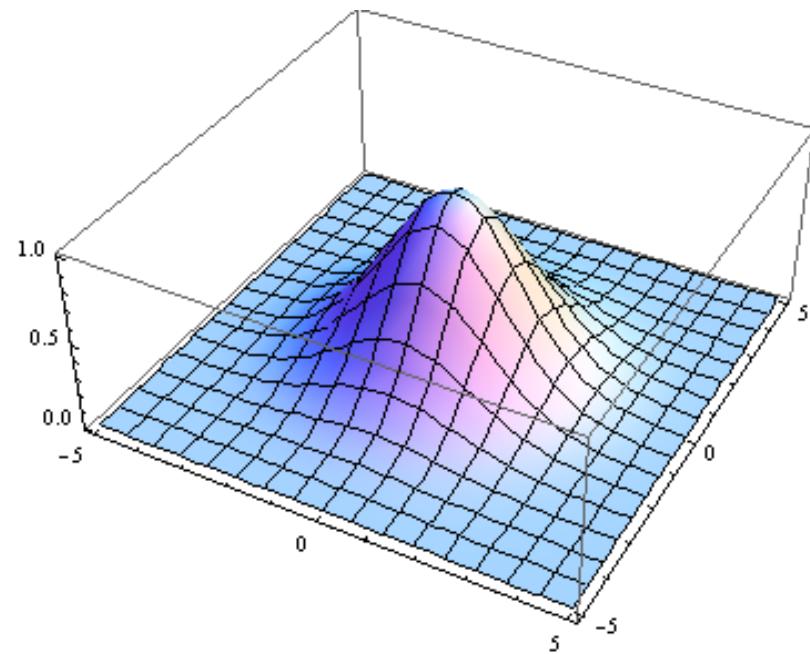


- Edges can be detected using:
 - The **maximum** of the first-order derivative
 - The **zero-crossing** of the second-order derivative

Gaussian Derivatives in 2D

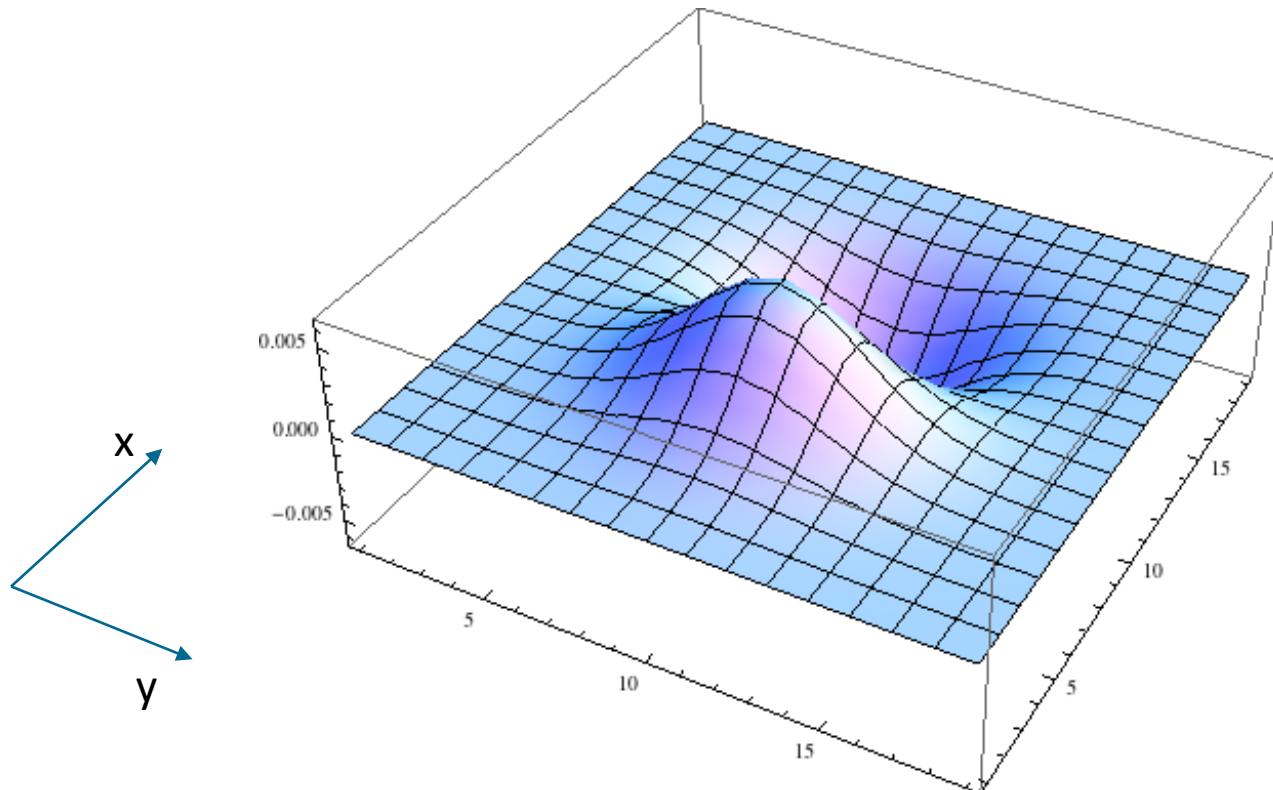
- 2D Gaussian kernel

$$G = e^{-\frac{x^2+y^2}{2\sigma^2}}$$



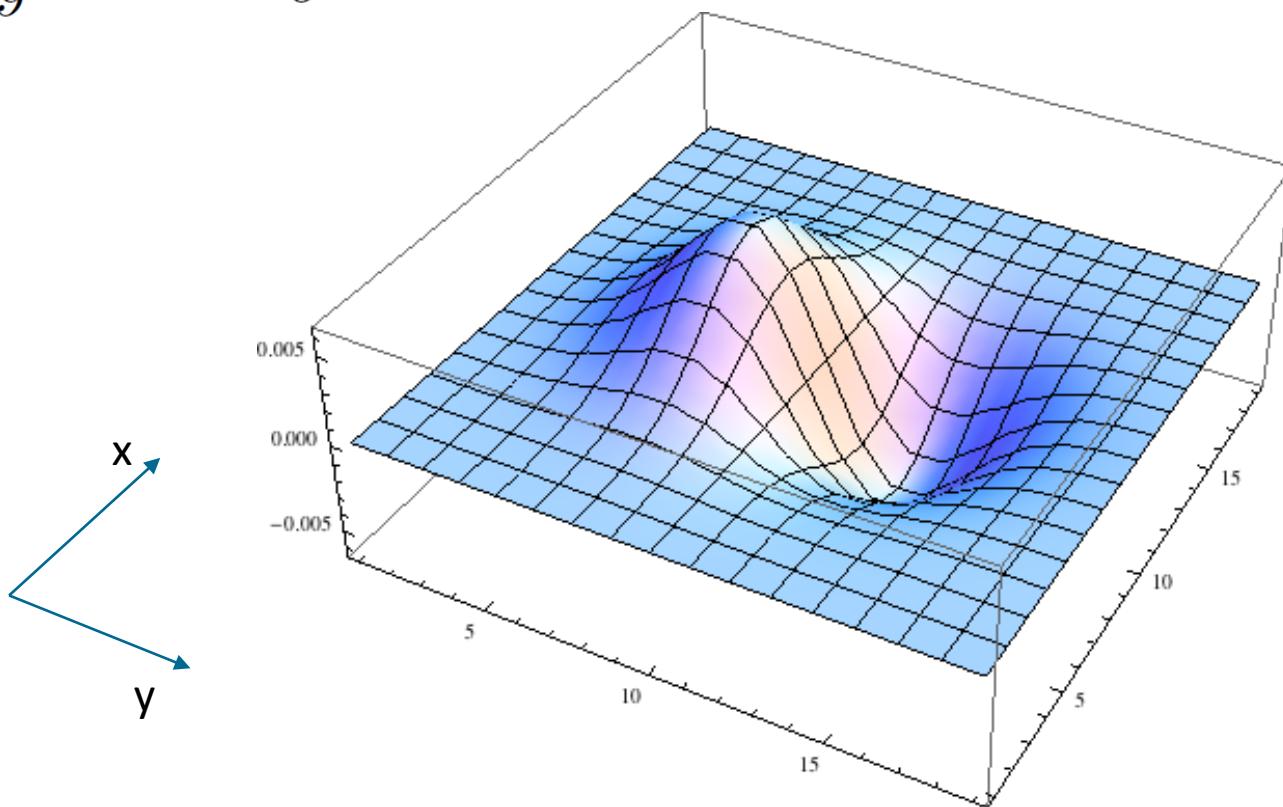
Gaussian Derivative

$$g'_x = \frac{\partial G(x, y)}{\partial x} = -\frac{x}{\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



Gaussian Derivative

$$g'_y = \frac{\partial G(x, y)}{\partial y} = -\frac{y}{\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



Edge detection – 1st order derivative

- Given the vector of the gradient [G_x G_y], we can compute the **magnitude** and the **direction** of the **gradient**:

$$\nabla \vec{f} = \begin{bmatrix} \frac{\partial f}{\partial x} & \frac{\partial f}{\partial y} \end{bmatrix} = \begin{bmatrix} G_x & G_y \end{bmatrix}$$

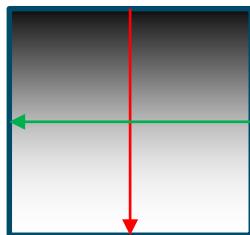
Magnitude $mag(\nabla \vec{f}) = \sqrt{G_x^2 + G_y^2}$

Direction $\alpha = \tan^{-1}\left(\frac{G_y}{G_x}\right)$

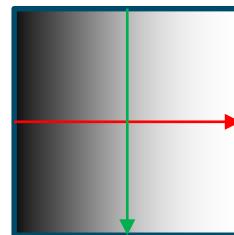
where α is the angle (in radians) from the x-axis to the point (x,y)

Edge detection - 1st order derivative

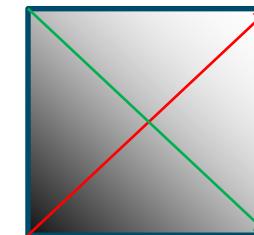
- The **gradient direction** gives the direction of maximum growth of the function from black to white



$$\left[0, \frac{\partial f}{\partial y} \right]$$



$$\left[\frac{\partial f}{\partial x}, 0 \right]$$

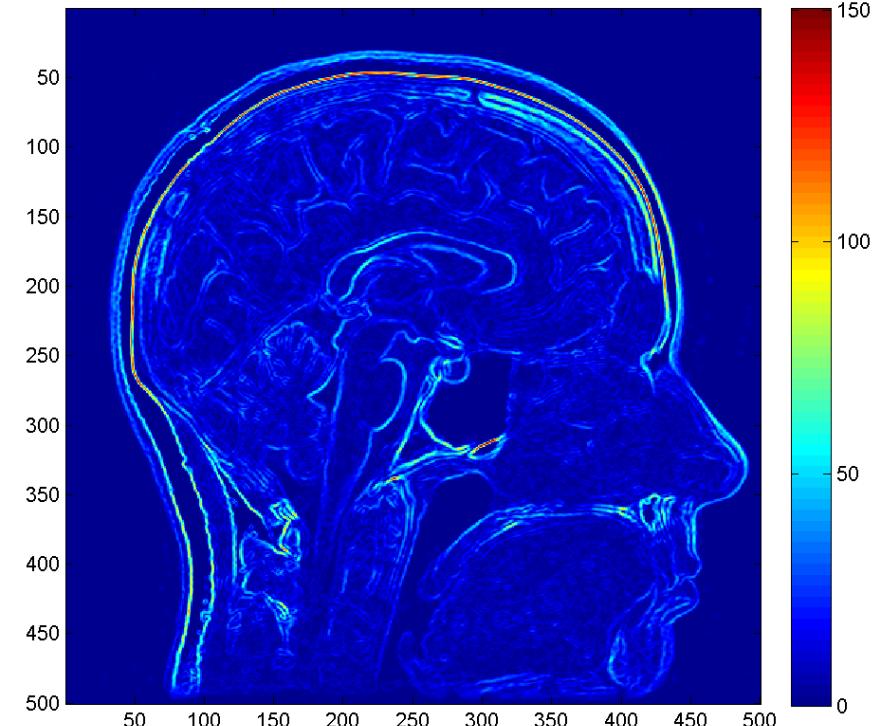
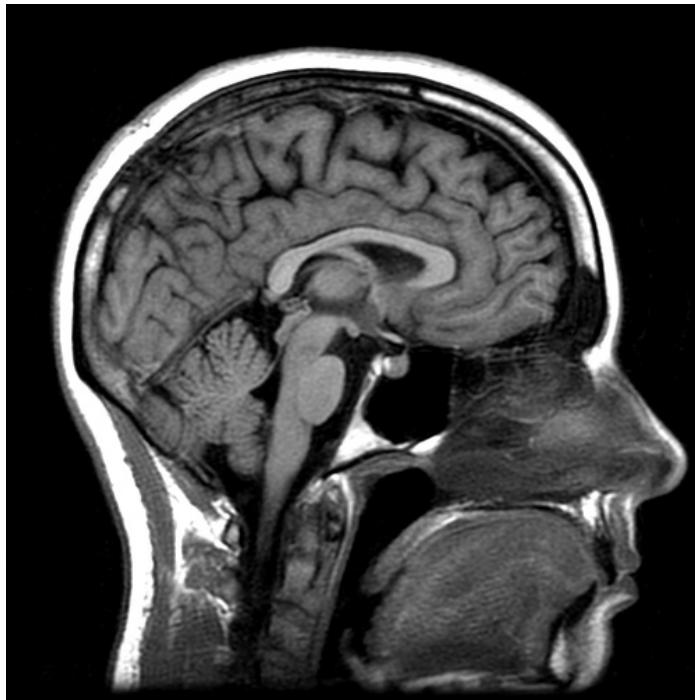


$$\left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

- The **edge magnitude** is the **magnitude** of the **gradient**
- The **edge direction** is rotated with respect to the gradient direction by -90°

Edge detection – 1st order derivative

$$mag(\vec{\nabla f}) = \sqrt{G_x^2 + G_y^2}$$



- Edges are the **local maxima** of the gradient magnitude
- We can classify **weak** and **strong** edges, based on the magnitude

In Python



SciPy.org



[Scipy.org](#)

[Docs](#)

[NumPy v1.12 Manual](#)

[NumPy Reference](#)

[Routines](#)

[Mathematical functions](#)

numpy.gradient

`numpy.gradient(f, *varargs, **kwargs)`

[\[source\]](#)

Return the gradient of an N-dimensional array.

The gradient is computed using second order accurate central differences in the interior and either first differences or second order accurate one-sides (forward or backwards) differences at the boundaries. The returned gradient hence has the same shape as the input array.

Parameters: `f : array_like`

An N-dimensional array containing samples of a scalar function.

`varargs : scalar or list of scalar, optional`

N scalars specifying the sample distances for each dimension, i.e. `dx, dy, dz, ...`. Default distance: 1. single scalar specifies sample distance for all dimensions. if `axis` is given, the number of varargs must equal the number of axes.

`edge_order : {1, 2}, optional`

Gradient is calculated using N-th order accurate differences at the boundaries. Default: 1.

New in version 1.9.1.

`axis : None or int or tuple of ints, optional`

Gradient is calculated only along the given axis or axes. The default (`axis = None`) is to calculate the gradient for all the axes of the input array. `axis` may be negative, in which case it counts from the last to the first axis.

New in version 1.11.0.

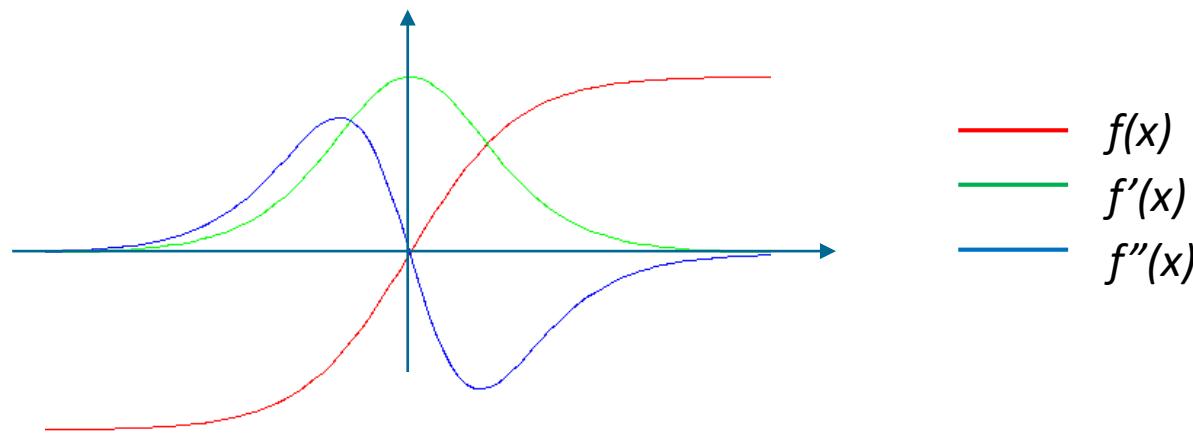
Returns:

`gradient : ndarray or list of ndarray`

A set of ndarrays (or a single ndarray if there is only one dimension) corresponding to the derivatives of `f` with respect to each dimension. Each derivative has the same shape as `f`.

Edge detection – 2nd order derivative

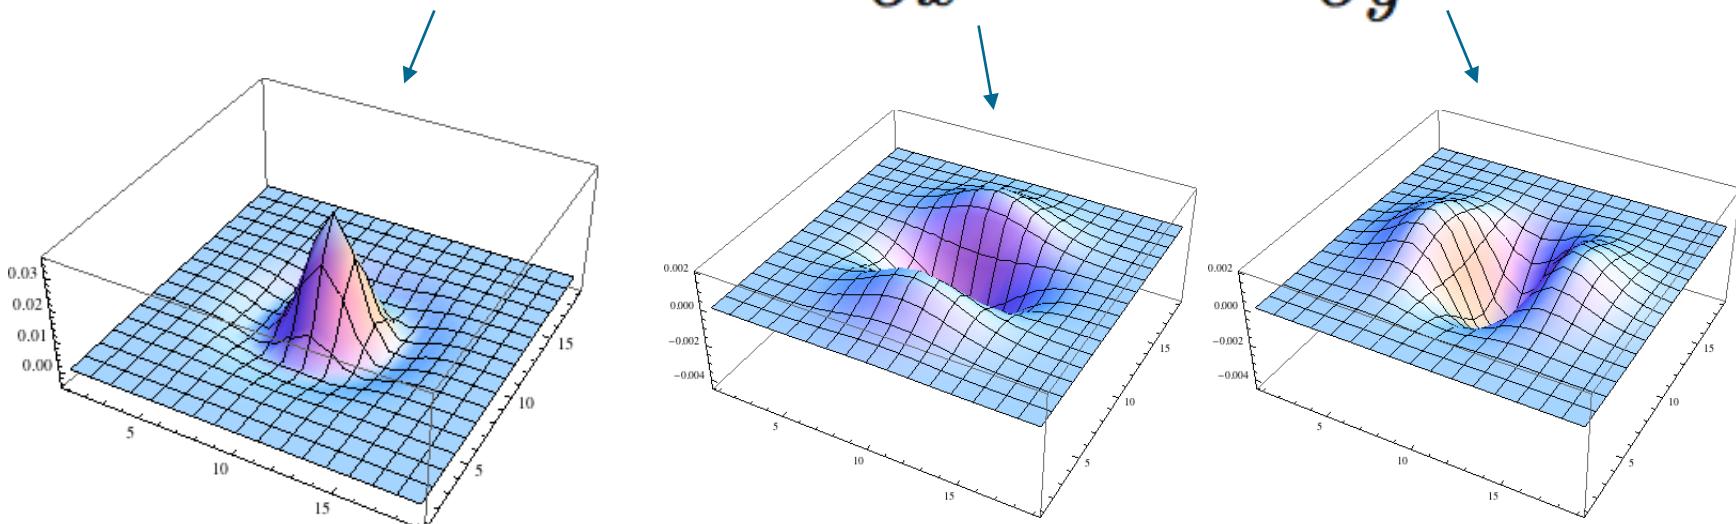
- In the presence of an edge, the second order derivative crosses the origin of the axis



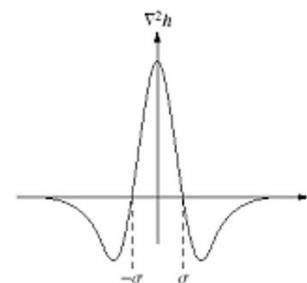
- The **zero-crossing** of the second-order derivative can be used to detect edges

Laplacian of Gaussian (LoG)

$$\nabla^2 G(x, y) = \frac{\partial^2 G(x, y)}{\partial x^2} + \frac{\partial^2 G(x, y)}{\partial y^2}$$

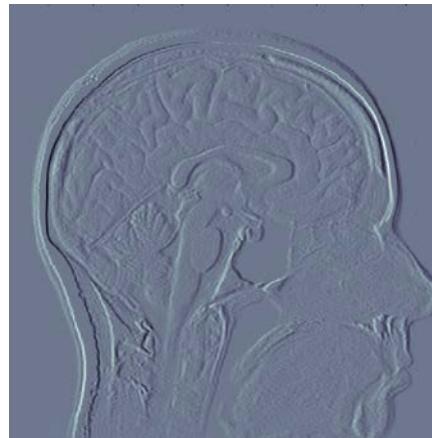


LoG a.k.a. "Mexican Hat"

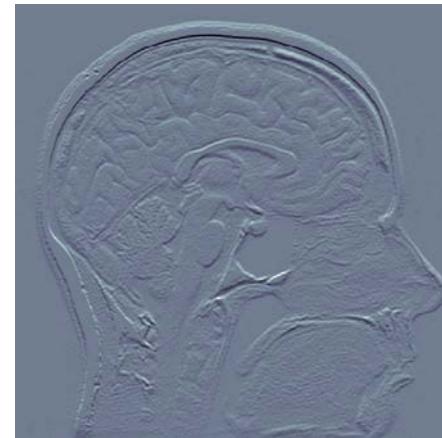


0	0	-1	0	0
0	-1	-2	-1	0
-1	-2	16	-2	-1
0	-1	-2	-1	0
0	0	-1	0	0

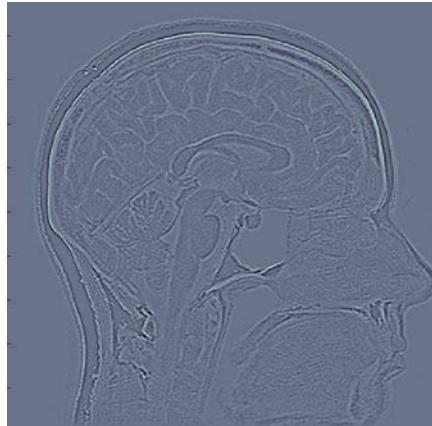
Gradient and Laplacian



Gx



Gy



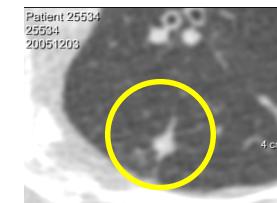
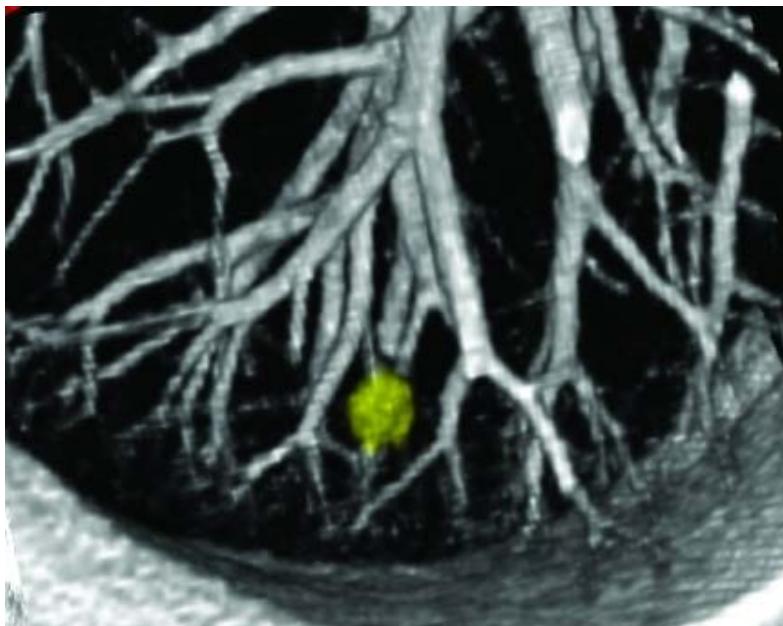
Δf

$$\Delta f = \nabla^2 f$$

Blob detector

- **Blob:** region of a digital image in which all the points can be considered to be similar to each other, and different from the surroundings

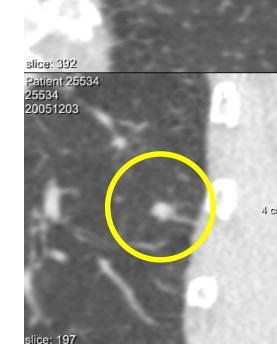
A pulmonary nodule fulfills the definition of blob, but in 3D...



axial



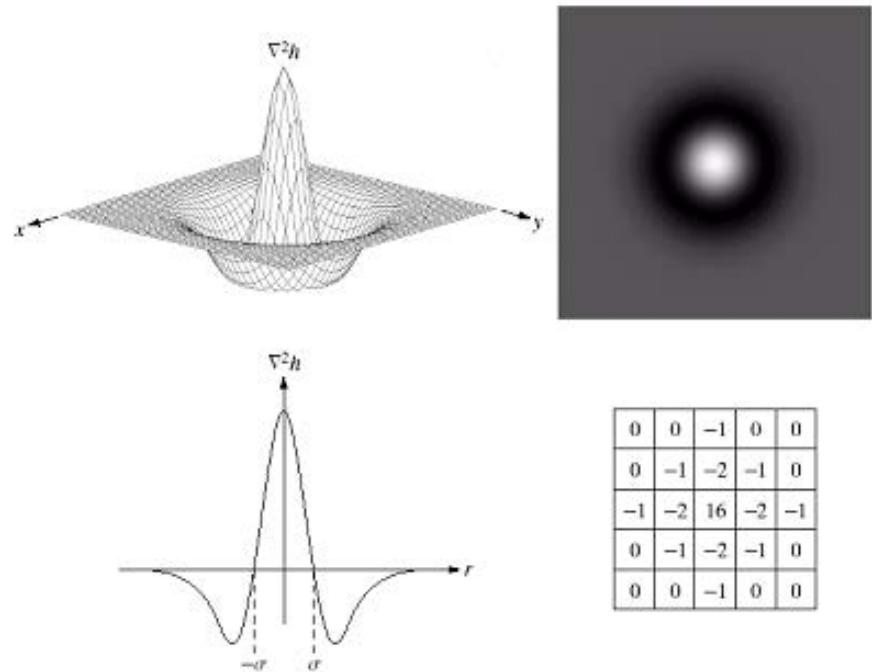
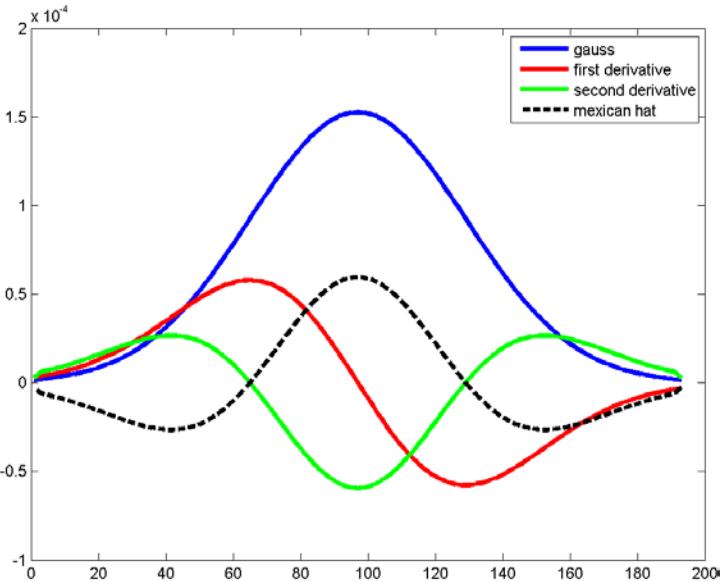
coronal



sagittal

Blob detector

- What's the response of the “Mexican hat” in the presence of a bright blob?



- Pulmonary nodules in 2D CT respond to a blob detector with the proper scale σ (related to the nodule size)

Scale space in practice

- How do I know what is the **optimal scale** for a given object of interest (e.g., pulmonary nodule)?
 - Consider a set of (relevant) sigma values -> filters
 - Apply all filters
 - Pick the one with the **maximum response**

Scale space in practice

- What's the optimal sigma value for a detection problem?
 - Consider a set of all relevant sigma values for a object of interest (e.g., pulmonary nodule)
 - Apply all filters independently
 - Combine the responses (sum / average)
 - You can assume that at least one of those will give a high response

Fast Fourier Transform

- (Spatial) convolution may be very **slow**, especially if **large** kernels are used
- We can work in the **frequency domain** using the Fourier transform

$$g(x, y) * h(x, y) \xrightarrow{F} G(u, v)H(u, v)$$

- In the frequency domain, a **convolution** becomes a **multiplication** of transformed image and kernel

In Python

[Scipy.org](#) [Docs](#) [SciPy v0.16.1 Reference Guide](#) [Signal processing \(scipy.signal\)](#)

scipy.signal.fftconvolve

`scipy.signal.fftconvolve(in1, in2, mode='full')`

Convolve two N-dimensional arrays using FFT.

Convolve `in1` and `in2` using the fast Fourier transform method, with the output size determined by the `mode` argument.

This is generally much faster than `convolve` for large arrays ($n > \sim 500$), but can be slower when only a few output values are needed, and can only output float arrays (int or object array inputs will be cast to float).

Parameters: `in1 : array_like`

First input.

`in2 : array_like`

Second input. Should have the same number of dimensions as `in1`; if sizes of `in1` and `in2` are not equal then `in1` has to be the larger array.

`mode : str {'full', 'valid', 'same'}, optional`

A string indicating the size of the output:

`full`

The output is the full discrete linear convolution of the inputs. (Default)

`valid`

The output consists only of those elements that do not rely on the zero-padding.

`same`

The output is the same size as `in1`, centered with respect to the 'full' output.

Returns:

`out : array`

An N-dimensional array containing a subset of the discrete linear convolution of `in1` with `in2`.

Paper for next week

Comparative study of retinal vessel segmentation methods on a new publicly available database

M. Niemeijer^{ab}, J. Staal^a, B. van Ginneken^a, M. Loog^a and M.D. Abràmoff^{abc}

^aImage Sciences Institute, Univ. Medical Center Utrecht, Utrecht, The Netherlands

^bDepartment of Ophthalmology, Vrije Univ. Medical Center, Amsterdam, The Netherlands

^cDepartment of Ophthalmology and Visual Sciences, Univ. of Iowa, USA

ABSTRACT

In this work we compare the performance of a number of vessel segmentation algorithms on a newly constructed retinal vessel image database. Retinal vessel segmentation is important for the detection of numerous eye diseases and plays an important role in automatic retinal disease screening systems. A large number of methods for retinal vessel segmentation have been published, yet an evaluation of these methods on a common database of screening images has not been performed. To compare the performance of retinal vessel segmentation methods we have constructed a large database of retinal images. The database contains forty images in which the vessel trees have been manually segmented. For twenty of those forty images a second independent manual segmentation is available. This allows for a comparison between the performance of automatic methods and the performance of a human observer. The database is available to the research community. Interested researchers are encouraged to upload their segmentation results to our website (<http://www.isi.uu.nl/Research/Databases>). The performance of five different algorithms has been compared. Four of these methods have been implemented as described in the literature. The fifth pixel classification based method was developed specifically for the segmentation of retinal vessels and is the only supervised method. We define the segmentation accuracy with respect to our gold standard as the performance measure. Results show that the pixel classification method performs best, but the second observer still performs significantly better.

Keywords: vessel segmentation, retina, comparative study, image database

Intelligent Systems in Medical Imaging

(NWI-IMC037)

Digital Image Processing

Francesco Ciompi

francesco.ciompi@radboudumc.nl

Radboud **umc**